



Alexis Balayre

High Performance Technical Computing Assignment

School of Aerospace, Transport and Manufacturing  
Computational Software of Techniques Engineering

MSc

Academic Year: 2023 - 2024

Supervisor: Dr Irene Moulitsas

5<sup>th</sup> February 2024

# Abstract

This paper explores the effectiveness of different parallelization strategies for multiplying sparse matrices by fat vectors, focusing on the use of MPI in High Performance Computing (HPC). It compares the performance of sequential and parallel methods, including row-by-row, column-by-column, and non-zero element approaches. The results highlight the implications of each strategy on execution time and efficiency, while considering the environmental impact of HPC, suggesting avenues towards more sustainable computing systems.

# Table of Contents

<b>Abstract</b>	<b>ii</b>
<b>Table of Contents</b>	<b>iii</b>
<b>List of Figures</b>	<b>vi</b>
<b>List of Tables</b>	<b>viii</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Methodology</b>	<b>2</b>
2.1 Problem Statement . . . . .	2
2.2 Data Structures . . . . .	3
2.2.1 Sparse Matrix . . . . .	3
2.2.2 Fat Vector . . . . .	3
2.3 Sequential Algorithm . . . . .	4
2.3.1 Algorithm Flow . . . . .	4
2.3.2 Temporal Complexity Analysis . . . . .	4
2.4 Line-Based Parallelism . . . . .	5
2.4.1 Algorithm Flow . . . . .	5
2.4.2 Time Complexity Analysis . . . . .	6
2.4.2.1 Computational Complexity . . . . .	6
2.4.2.2 Communication Complexity . . . . .	6
2.4.2.3 Final Result Reconstruction . . . . .	6
2.4.2.4 Overall Time Complexity . . . . .	7
2.4.3 Performance Analysis . . . . .	7
2.4.3.1 Performance Dependence . . . . .	7
2.4.3.2 Expected Performance . . . . .	7
2.5 Column-Wise Parallelism . . . . .	8
2.5.1 Algorithm Flow . . . . .	8
2.5.2 Temporal Complexity Analysis . . . . .	9
2.5.2.1 Computation Time Complexity . . . . .	9
2.5.2.2 Communication Time Complexity . . . . .	9
2.5.2.3 Final Result Reconstruction . . . . .	9
2.5.2.4 Overall Time Complexity . . . . .	9
2.5.3 Performance Analysis . . . . .	10
2.5.3.1 Performance Dependence . . . . .	10

2.5.3.2	Expected Performance . . . . .	10
2.6	Non-Zero Element Parallelism . . . . .	11
2.6.1	Algorithm Flow . . . . .	11
2.6.2	Temporal Complexity Analysis . . . . .	12
2.6.3	Computation Complexity . . . . .	12
2.6.4	Communication Complexity . . . . .	12
2.6.4.1	Final Result Reconstruction . . . . .	12
2.6.5	Overall Time Complexity . . . . .	12
2.6.6	Performance Analysis . . . . .	13
2.6.6.1	Performance Dependence . . . . .	13
2.6.6.2	Expected Performance . . . . .	13
2.7	Performance Metrics . . . . .	14
<b>3</b>	<b>Results and Discussion</b>	<b>16</b>
3.1	Results . . . . .	16
3.1.1	Sparse Matrix Impact . . . . .	16
3.1.1.1	Execution Time Evolution . . . . .	17
3.1.1.2	Average Communication Time Evolution . . . . .	19
3.1.1.3	Average Computation Time Evolution . . . . .	21
3.1.1.4	Performance Evolution . . . . .	23
3.1.1.5	Performance Summary . . . . .	25
3.1.2	Fat Vector Impact . . . . .	26
3.1.2.1	Execution Time Evolution . . . . .	26
3.1.2.2	Average Communication Time Evolution . . . . .	28
3.1.2.3	Average Computation Time Evolution . . . . .	30
3.1.2.4	Performance Evolution . . . . .	32
3.1.2.5	Performance Summary . . . . .	34
3.2	HPC's Environmental Impact . . . . .	35
<b>4</b>	<b>Conclusion</b>	<b>36</b>
	<b>References</b>	<b>37</b>
<b>A</b>	<b>Documentation</b>	<b>38</b>
A.A	Project tree . . . . .	38
A.B	Getting Started . . . . .	38
A.C	Methods Overview . . . . .	39
A.C.1	Utils.h . . . . .	39
A.C.1.1	ConvertPETScMatToFatVector . . . . .	39
A.C.1.2	areMatricesEqual . . . . .	39
A.C.1.3	readMatrixMarketFile . . . . .	39
A.C.1.4	generateLargeFatVector . . . . .	40
A.C.1.5	serialize and deserialize . . . . .	40
A.C.2	SparseMatrixFatVectorMultiply.h . . . . .	40
A.C.2.1	sparseMatrixFatVectorMultiply . . . . .	40
A.C.3	SparseMatrixFatVectorMultiplyRowWise.h . . . . .	41
A.C.3.1	sparseMatrixFatVectorMultiplyRowWise . . . . .	41

A.C.4	SparseMatrixFatVectorMultiplyColumnWise.h . . . . .	41
A.C.4.1	sparseMatrixFatVectorMultiplyColumnWise . . . . .	41
A.C.5	SparseMatrixFatVectorMultiplyNonZeroElement.h . . . . .	41
A.C.5.1	sparseMatrixFatVectorMultiplyNonZeroElement . . . . .	41
<b>B</b>	<b>Source Codes</b>	<b>42</b>
B.A	Data Structures . . . . .	42
B.B	Sequential Algorithm . . . . .	43
B.B.1	Declaration File . . . . .	43
B.B.2	Implementation File . . . . .	43
B.C	Line-Based Parallelism . . . . .	44
B.C.1	Declaration File . . . . .	44
B.C.2	Implementation File . . . . .	44
B.D	Column-Wise Parallelism . . . . .	47
B.D.1	Declaration File . . . . .	47
B.D.2	Implementation File . . . . .	47
B.E	Non-Zero Element Parallelism . . . . .	50
B.E.1	Declaration File . . . . .	50
B.E.2	Implementation File . . . . .	50
B.F	Utility Functions . . . . .	53
B.F.1	Declaration File . . . . .	53
B.F.2	Implementation File . . . . .	54
B.G	Main File . . . . .	58
B.H	Scripts . . . . .	65
B.H.1	MPI Submission Script . . . . .	65
B.H.2	Batch Test Script . . . . .	67
B.H.3	Get CSV Script . . . . .	69

# List of Figures

3.1	Cage4 Matrix Execution Time Evolution . . . . .	17
3.2	DC1 Matrix Execution Time Evolution . . . . .	17
3.3	Amazon0302 Matrix Execution Time Evolution . . . . .	18
3.4	Cop20k_A Matrix Execution Time Evolution . . . . .	18
3.5	FEM_3D_thermal1 Matrix Execution Time Evolution . . . . .	18
3.6	Cage4 Matrix Communication Time Evolution . . . . .	19
3.7	DC1 Matrix Communication Time Evolution . . . . .	19
3.8	Amazon0302 Matrix Communication Time Evolution . . . . .	20
3.9	Cop20k_A Matrix Communication Time Evolution . . . . .	20
3.10	FEM_3D_thermal1 Matrix Communication Time Evolution . . . . .	20
3.11	Cage4 Matrix Computation Time Evolution . . . . .	21
3.12	DC1 Matrix Computation Time Evolution . . . . .	21
3.13	Amazon0302 Matrix Computation Time Evolution . . . . .	22
3.14	Cop20k_A Matrix Computation Time Evolution . . . . .	22
3.15	FEM_3D_thermal1 Matrix Computation Time Evolution . . . . .	22
3.16	Cage4 Matrix Performance Evolution . . . . .	23
3.17	DC1 Matrix Performance Evolution . . . . .	23
3.18	Amazon0302 Matrix Performance Evolution . . . . .	24
3.19	Cop20k_A Matrix Performance Evolution . . . . .	24
3.20	FEM_3D_thermal1 Matrix Performance Evolution . . . . .	24
3.21	Execution Time Evolution (k=1) . . . . .	26
3.22	Execution Time Evolution (k=3) . . . . .	26
3.23	Execution Time Evolution (k=6) . . . . .	27
3.24	Execution Time Evolution (k=9) . . . . .	27
3.25	Execution Time Evolution (k=12) . . . . .	27
3.26	Communication Time Evolution (k=1) . . . . .	28
3.27	Communication Time Evolution (k=3) . . . . .	28
3.28	Communication Time Evolution (k=6) . . . . .	29
3.29	Communication Time Evolution (k=9) . . . . .	29
3.30	Communication Time Evolution (k=12) . . . . .	29
3.31	Computation Time Evolution (k=1) . . . . .	30
3.32	Computation Time Evolution (k=3) . . . . .	30
3.33	Computation Time Evolution (k=6) . . . . .	31
3.34	Computation Time Evolution (k=9) . . . . .	31
3.35	Computation Time Evolution (k=12) . . . . .	31
3.36	Performance Evolution (k=1) . . . . .	32

3.37	Performance Evolution (k=3) . . . . .	32
3.38	Performance Evolution (k=6) . . . . .	33
3.39	Performance Evolution (k=9) . . . . .	33
3.40	Performance Evolution (k=12) . . . . .	33

# List of Tables

3.1	Sparse matrix specifications . . . . .	16
-----	--	----



# Chapter 1

## Introduction

High-Performance Computing (HPC) is a branch of computing that uses supercomputers and server clusters to solve complex, computationally intensive problems. Unlike a personal computer with a single processor, an HPC system is made up of many processors working in parallel, considerably increasing processing capacity. This enables scientists and engineers to carry out detailed numerical simulations, such as forecasting the weather or solving structural engineering problems.

Cranfield University has two HPC systems: CRESCENT2 and DELTA. However, this report will focus exclusively on CRESCENT2, which is an HPC cluster designed to provide computing power for teaching and research. CRESCENT 2 nodes are equipped with Intel Xeon E5 2620 processors, and each node contains two 16-core processors and 16 gigabytes of RAM.

The aim of this report is to explore distributed memory parallel programming strategies for optimising the performance of sparse matrix multiplication by a fat vector, a common operation in numerical linear algebra.

# Chapter 2

## Methodology

### 2.1 Problem Statement

Consider a sparse matrix  $M$  of dimensions  $m \times n$  and a fat vector  $v$  of dimensions  $n \times k$ . The objective is to perform the multiplication  $M \times v$ , yielding a result that is of dimensions  $m \times k$ .

The matrix  $M$  is defined as:

$$M = \begin{pmatrix} m_{11} & m_{12} & \cdots & m_{1n} \\ m_{21} & m_{22} & \cdots & m_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ m_{m1} & m_{m2} & \cdots & m_{mn} \end{pmatrix} \quad (2.1)$$

where most elements of  $M$  are zeros.

The vector  $v$  is defined as:

$$v = \begin{pmatrix} v_{11} & v_{12} & \cdots & v_{1k} \\ v_{21} & v_{22} & \cdots & v_{2k} \\ \vdots & \vdots & \ddots & \vdots \\ v_{n1} & v_{n2} & \cdots & v_{nk} \end{pmatrix} \quad (2.2)$$

## 2.2 Data Structures

In numerical computation and linear algebra, efficient use of memory and fast computation are crucial. This is particularly true when working with hollow matrices and fat vectors.

### 2.2.1 Sparse Matrix

The sparse matrix is represented in CSR (Compressed Sparse Row) format, which is particularly effective for storing and manipulating matrices where the majority of elements are zero. The CSR structure consists of three main vectors:

- **values**: A vector storing all the non-zero elements of the matrix.
- **rowPtr**: A vector storing the starting index for each element in the *values* vector.
- **colIndices**: A vector storing the column indices for each element in the vector *values*.

Here is an example of a sparse matrix in CSR format:

- `values = {1, 2, 3, 4}`
- `rowPtr = {0, 2, 3, 3, 4}`
- `colIndices = {0, 2, 2, 3}`

This hollow matrix can be visualised as:

$$\begin{bmatrix} 1 & 0 & 2 & 0 \\ 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 4 \end{bmatrix}$$

The **SparseMatrix** structure is defined in Appendix B.A.

### 2.2.2 Fat Vector

Unlike a hollow matrix, a fat vector (illustrated by equation 2.2) stores all its elements, including zeros. The data structure for a fat vector is a two-dimensional array, where each row represents a separate vector. The **FatVector** structure is defined in Appendix B.A.

## 2.3 Sequential Algorithm

Let  $M$  be a sparse matrix of size  $m \times n$  with  $z$  non-zero elements, stored in CSR format, and  $v$  be a fat vector of size  $n \times k$ . The sequential algorithm for multiplying  $M$  by  $v$  is implemented in Appendix B.B.

### 2.3.1 Algorithm Flow

---

**Algorithm 1** Sequential algorithm

---

**Require:**  $M$  is an  $m \times n$  sparse matrix

**Require:**  $v$  is an  $n \times k$  fat vector

**Ensure:**  $Result$  is an  $m \times k$  matrix

$Result \leftarrow$  zero matrix of size  $m \times k$

**for**  $i \leftarrow 0$  **to**  $m - 1$  **do**

**for** each non-zero element  $(j, \text{value})$  in row  $i$  of  $M$  **do**

**for**  $j \leftarrow 0$  **to**  $k - 1$  **do**

$Result[i][l] \leftarrow Result[i][l] + (\text{value} \times v[j][l])$

**end for**

**end for**

**end for**

**return**  $Result$

---

The algorithmic flow can be more explicitly detailed by:

1. **Initialisation:** Create a zero matrix of size  $m \times k$  to store the result.
2. **Row-wise Processing:** Iterate over each row  $i$  of matrix  $M$ , leveraging the CSR format to efficiently access non-zero elements.
3. **Element-wise Multiplication and Accumulation:** For each non-zero element in row  $i$ , identified by its column index  $j$  and value, conduct a nested iteration over the columns  $l$  of vector  $v$ , multiplying the non-zero element by the corresponding vector element and accumulating the result in  $Result[i][l]$ .

### 2.3.2 Temporal Complexity Analysis

Given a sparse matrix  $M$  of size  $m \times n$  with  $z$  non-zero elements and a fat vector  $v$  of size  $n \times k$ , the serial algorithm for multiplying  $M \times v$  iterates through each non-zero element of the matrix  $M$  to compute the product.

The algorithm performs two operations (a multiplication and an addition) for each non-zero element with respect to each column of  $v$ , resulting in a total of  $2zk$  operations.

Hence, the time complexity of the serial sparse matrix-fat vector multiplication algorithm can be expressed as:

$$T(n) = O(zk) \tag{2.3}$$

## 2.4 Line-Based Parallelism

This algorithm distributes the rows of the sparse matrix across multiple processes for parallel computation in a line-based manner. The implementation is detailed in Appendix B.C.

### 2.4.1 Algorithm Flow

---

**Algorithm 2** Row-wise Parallel Sparse Matrix-Fat Vector Multiplication

---

**Require:**  $M$  is an  $m \times n$  sparse matrix  
**Require:**  $v$  is an  $n \times k$  fat vector  
**Require:**  $worldSize$  is the number of processes  
**Require:**  $worldRank$  is the rank of the current process  
**Ensure:**  $finalResult$  is an  $m \times k$  matrix, result of  $M \times v$

```

rowsPerProcess  $\leftarrow m / worldSize$ 
extraRows  $\leftarrow m \bmod worldSize$ 
startRow  $\leftarrow worldRank \times rowsPerProcess + \min(worldRank, extraRows)$ 
endRow  $\leftarrow startRow + rowsPerProcess$ 
if  $worldRank < extraRows$  then
    endRow  $\leftarrow endRow + 1$ 
end if
Initialise  $localResult$  with zeros of size  $(endRow - startRow) \times k$ 
for  $i \leftarrow startRow$  to  $endRow - 1$  do
    for each non-zero element  $(j, value)$  in row  $i$  of  $M$  do
        for  $k \leftarrow 0$  to  $k - 1$  do
            localIndex  $\leftarrow (i - startRow) \times k + k$ 
            localResult[localIndex]  $\leftarrow localResult[localIndex] + value \times v[j][k]$ 
        end for
    end for
end for
if  $worldRank == 0$  then
    Prepare  $recvCounts$  and  $displacements$  for gathering
end if
MPI_Gatherv(localResult, ...)
if  $worldRank == 0$  then
    Reassemble  $finalResult$  from all  $localResults$ 
    return  $finalResult$ 
end if

```

---

The algorithmic flow can be explicitly detailed by:

1. **Initialisation:** Obtain MPI world size and rank to determine each process's role.
2. **Row Distribution:** Assign a subset of rows from the sparse matrix to each process based on the rank, ensuring an even distribution with possible adjustments for any remainder.

3. **Local Computation:** Each process calculates the product of its assigned rows with the fat vector, storing results in a local vector.
4. **Gather Results:** Use `MPI_Gatherv` to collect the local result vectors from all processes into a single vector at the root process.
5. **Final Result Reconstruction:** The root process reconstructs the final result matrix from the gathered vector.

## 2.4.2 Time Complexity Analysis

Given a sparse matrix  $M$  of size  $m \times n$  with  $z$  non-zero elements and a fat vector  $v$  of size  $n \times k$ , the line-based algorithm distributes the multiplication task across  $p$  processors.

### 2.4.2.1 Computational Complexity

Each process computes a portion of the result vector, working on approximately  $m/p$  rows of the sparse matrix. The computation time ( $T_{\text{comp}}$ ) is therefore influenced by the distribution of non-zero elements across the rows. Assuming a uniform distribution, the computation time can be estimated as:

$$T_{\text{comp}} = O\left(\frac{z \times k}{p}\right) \quad (2.4)$$

### 2.4.2.2 Communication Complexity

After computation, each process holds a part of the result vector that needs to be combined to form the final result. The main communication cost arises from gathering these parts at a single process or distributing them among all processes.

- **Startup Overhead:** The initiation of a communication incurs a latency cost,  $\alpha$ , which is significant when the number of messages is large.
- **Data Transmission Cost:** Each process sends its portion of the result vector, incurring a cost proportional to the amount of data sent, the number of processes and the network bandwidth, denoted by  $\beta$ .

The communication time complexity ( $T_{\text{comm}}$ ) can be approximated as:

$$T_{\text{comm}} = p \times \alpha + \beta \times \frac{m}{p} \times k \quad (2.5)$$

### 2.4.2.3 Final Result Reconstruction

After gathering the local results, the root process reconstructs the final result matrix. This step has a time complexity proportional to the size of the final matrix,  $O(m \times k)$ .

### 2.4.2.4 Overall Time Complexity

The overall time complexity, including both computation and communication, is given by:

$$T_{\text{overall}} = T_{\text{comp}} + T_{\text{comm}} = O\left(\frac{z \times k}{p}\right) + p \times \alpha + \beta \times \frac{m}{p} \times k + O(m \times k) \quad (2.6)$$

## 2.4.3 Performance Analysis

### 2.4.3.1 Performance Dependence

- Number of processes ( $p$ ): Performance improves with  $p$  until communication overhead becomes significant.
- Number of rows ( $m$ ): Affects workload distribution. Performance optimal when  $p \leq m$ .
- Number of non-zero elements ( $z$ ): Higher  $z$  increases computation workload but mitigated by parallel execution.
- Number of columns ( $k$ ): Increases computation linearly. Each row computation involves all columns.

### 2.4.3.2 Expected Performance

Optimal when  $p \leq m$  with even workload distribution. Performance may degrade if  $p > m$  due to idle processes or when communication overhead outweighs computation benefits.

## 2.5 Column-Wise Parallelism

This algorithm distributes the columns of the fat vector across multiple processes for parallel computation in a column-based manner. The implementation is detailed in Appendix B.D.

### 2.5.1 Algorithm Flow

---

**Algorithm 3** Column-wise Parallel Sparse Matrix-Fat Vector Multiplication

---

**Require:**  $M$  is an  $m \times n$  sparse matrix  
**Require:**  $v$  is an  $n \times k$  fat vector  
**Require:**  $worldSize$  is the number of processes  
**Require:**  $worldRank$  is the rank of the current process  
**Ensure:**  $finalResult$  is an  $m \times k$  matrix, result of  $M \times v$

```

 $colsPerProcess \leftarrow k / worldSize$ 
 $extraCols \leftarrow k \bmod worldSize$ 
 $startCol \leftarrow worldRank \times colsPerProcess$ 
 $endCol \leftarrow startCol + colsPerProcess$ 
if  $worldRank == worldSize - 1$  then
     $endCol \leftarrow endCol + extraCols$ 
end if
 $localSize \leftarrow m \times (endCol - startCol)$ 
Initialise  $localResult$  with zeros of size  $localSize$ 
for  $col \leftarrow startCol$  to  $endCol - 1$  do
    for  $row \leftarrow 0$  to  $m - 1$  do
         $sum \leftarrow 0$ 
        for each non-zero element  $(i, value)$  in row  $row$  of  $M$  do
             $sum \leftarrow sum + value \times v[i][col]$ 
        end for
         $localResult[row][col - startCol] \leftarrow sum$ 
    end for
end for
if  $worldRank == 0$  then
    Initialise  $finalResult$  with zeros of size  $m \times k$ 
end if
Gather  $localResult$  from all processes to  $finalResult$  at root
if  $worldRank == 0$  then
    State Reassemble  $finalResult$  from gathered  $localResults$ 
    return  $finalResult$ 
end if

```

---

The algorithmic flow can be explicitly detailed by:

1. **Initialisation:** Obtain MPI world size and rank to determine each process's role.



2. **Column Distribution:** Calculate the number of columns each process will handle, distributing any extra columns to the last processes, and define the start and end column indices for each process.
3. **Local Computation:** Each process computes a portion of the multiplication result for its assigned columns, iterating through the sparse matrix rows and the relevant columns of the fat vector.
4. **Gather Results:** Use `MPI_Gatherv` to collect the local results from all processes into a single result vector at the root process.
5. **Final Result Reconstruction:** The root process reassembles the gathered results into the final fat vector matrix, ensuring the elements are correctly positioned according to their original indices.

## 2.5.2 Temporal Complexity Analysis

Consider a sparse matrix  $M$  of size  $m \times n$  with  $z$  non-zero elements and a fat vector  $v$  of size  $n \times k$ , the column-wise algorithm distributes the multiplication task across  $p$  processors.

### 2.5.2.1 Computation Time Complexity

Each process is responsible for approximately  $k/p$  columns of the fat vector. The computation time ( $T_{\text{comp}}$ ) is therefore influenced by the distribution of non-zero elements across the columns. Assuming a uniform distribution, the computation time can be estimated as:

$$T_{\text{comp}} = O\left(\frac{k \times z}{p}\right) \quad (2.7)$$

### 2.5.2.2 Communication Time Complexity

After local computations, partial results must be gathered at the root process. The communication time complexity ( $T_{\text{comm}}$ ) can be approximated as:

$$T_{\text{comm}} = p \times \alpha + \beta \times m \times \frac{k}{p} \quad (2.8)$$

### 2.5.2.3 Final Result Reconstruction

After gathering the local results, the root process reconstructs the final result matrix. This step has a time complexity proportional to the size of the final matrix,  $O(m \times k)$ .

### 2.5.2.4 Overall Time Complexity

The overall time complexity of the column-wise parallel sparse matrix-vector multiplication algorithm is dominated by the sum of computation and communication complexities:

$$T_{\text{overall}} = T_{\text{comp}} + T_{\text{comm}} = O\left(\frac{k \times z}{p}\right) + p \times \alpha + \beta \times m \times \frac{k}{p} + O(m \times k) \quad (2.9)$$

### 2.5.3 Performance Analysis

#### 2.5.3.1 Performance Dependence

- Number of processes ( $p$ ): Performance improvement depends on the ratio of  $k$  to  $p$ .
- Number of rows ( $m$ ): Less impact on performance scaling compared to  $p$  and  $k$ .
- Number of non-zero elements ( $z$ ): Impacts computation time, balanced by parallel processing.
- Number of columns ( $k$ ): Critical for performance. Optimal when large and divisible by  $p$ .

#### 2.5.3.2 Expected Performance

Best when  $k$  is significantly larger than  $p$  and divisible, allowing efficient parallel processing with minimal communication overhead. Performance may degrade if  $p > k$  due to idle processes or when communication outweighs computation benefits.

## 2.6 Non-Zero Element Parallelism

This algorithm distributes the non-zero elements of the sparse matrix across multiple processes for parallel computation. The implementation is detailed in Appendix B.E.

### 2.6.1 Algorithm Flow

---

**Algorithm 4** Non-Zero Element Parallel Sparse Matrix-Fat Vector Multiplication

---

**Require:**  $M$  is an  $m \times n$  sparse matrix

**Require:**  $v$  is an  $n \times k$  fat vector

**Require:**  $worldSize$  is the number of processes

**Require:**  $worldRank$  is the rank of the current process

**Ensure:**  $finalResult$  is an  $m \times k$  matrix, result of  $M \times v$

Calculate the total number of non-zero elements and distribute them among MPI processes

Determine  $startIdx$  and  $endIdx$  for non-zero elements for the current process

Map non-zero element indices to their corresponding row indices in the sparse matrix

Initialise  $localResult$  with zeros of size  $m \times k$

**for**  $idx \leftarrow startIdx$  **to**  $endIdx - 1$  **do**

    Determine  $row$ ,  $col$ , and  $value$  for each non-zero element

**for**  $k \leftarrow 0$  **to**  $k - 1$  **do**

$localResult[row \times k + k] \leftarrow localResult[row \times k + k] + value \times v[col][k]$

**end for**

**end for**

Use MPI\_Reduce to sum up  $localResults$  from all processes to  $flatFinalResult$  at the root process

**if**  $worldRank == 0$  **then**

    Reconstruct  $finalResult$  from  $flatFinalResult$

**return**  $finalResult$

**end if**

---

The algorithmic flow can be explicitly detailed by:

1. **Initialisation:** Obtain MPI world size and rank to determine each process's role.
2. **Non-Zero Elements Distribution:** Calculate each process's share of non-zero elements in the sparse matrix.
3. **Local Computation:** Each process multiplies its assigned non-zero elements with corresponding columns in the fat vector, accumulating results locally.
4. **Gather Results:** Use Reduce to sum up all local results into a single vector on the root process.
5. **Final Result Reconstruction:** The root process reconstructs the final result matrix from the gathered vector.

### 2.6.2 Temporal Complexity Analysis

Given a sparse matrix  $M$  of size  $m \times n$  with  $z$  non-zero elements and a fat vector  $v$  of size  $n \times k$ , the non-zero elements algorithm distributes the multiplication task across  $p$  processors.

### 2.6.3 Computation Complexity

Each process is responsible for a subset of the non-zero elements. The total computation workload is proportional to the number of non-zero elements  $z$ , distributed evenly among  $p$  processes:

$$T_{\text{comp}} = O\left(\frac{z \times k}{p}\right) \quad (2.10)$$

Given that each non-zero element computation involves a multiplication and an addition, the computation complexity remains linear with respect to the number of non-zero elements handled by each process.

### 2.6.4 Communication Complexity

The communication complexity involves reducing the local results from all the processes to a single result at the root process. The communication time complexity ( $T_{\text{comm}}$ ) can be approximated as:

$$T_{\text{comm}} = p \times \alpha + \beta \times m \times \frac{k}{p} \quad (2.11)$$

where:

- $\alpha$  is the latency cost of initiating a communication.
- $\beta$  is the cost of data transmission.

#### 2.6.4.1 Final Result Reconstruction

After reducing the local results, the root process reconstructs the final result matrix. This step has a time complexity proportional to the size of the final matrix,  $O(m \times k)$ .

### 2.6.5 Overall Time Complexity

The overall time complexity of the SparseMatrixFatVectorMultiplyNonZeroElement algorithm is the sum of computation and communication complexities:

$$T_{\text{overall}} = T_{\text{comp}} + T_{\text{comm}} = O\left(\frac{z \times k}{p}\right) + p \times \alpha + \beta \times m \times \frac{k}{p} + O(m \times k) \quad (2.12)$$

## 2.6.6 Performance Analysis

### 2.6.6.1 Performance Dependence

- Number of processes ( $p$ ): Performance improves with  $p$ , but aggregation communication can be a bottleneck.
- Number of rows ( $m$ ): Impacts performance through non-zero element distribution.
- Number of non-zero elements ( $z$ ): Directly proportional to computation workload, benefiting from parallel execution.
- Number of columns ( $k$ ): Increases workload linearly, balanced by distributing non-zero elements across processes.

### 2.6.6.2 Expected Performance

Highly efficient for matrices with large and evenly distributed  $z$ , maximising parallelism. Optimal when  $z/p$  is large, minimising idle time and communication overhead.

## 2.7 Performance Metrics

The evaluation of algorithm performance within the main program (detailed in appendix B.G) involves a systematic approach to assess efficiency across various implementations. In addition, the PETSc library (1) is used to compare the performance of the parallel algorithms with a highly optimised parallel sparse matrix-vector multiplication implementation. The program’s workflow is structured as follows:

1. **Initialisation:** Initial setup of MPI and PETSc environments to enable distributed computation and matrix operations.
2. **Matrix and Vector Preparation:**
  - *Matrix Reading:* Sparse matrices are sourced from files utilising the `readMatrixMarketFile` method (refer to appendix B.F), leveraging the Matrix Market I/O library (2) for efficient data handling.
  - *Fat Vector Generation:* Creation of fat vectors, with dimensions tailored to the corresponding matrices and predetermined column counts, ensuring compatibility for multiplication.
  - *Distribution:* Uniform distribution of matrices and vectors across processes for parallel computation.
3. **Serial Multiplication:** Execution and timing of matrix-vector multiplication in a serial context to establish a performance baseline.
4. **Parallel Multiplication Variants:** Application of distinct parallel multiplication strategies—line-based, column-based, and non-zero element—to measure execution times and identify efficiency variances.
5. **PETSc Implementation:**
  - *Conversion:* Adaptation of matrices and vectors to PETSc formats to utilise its optimized operations.
  - *Execution:* Conducting matrix-vector multiplication within the PETSc framework.
  - *Result Conversion:* Reformatting PETSc output back to Fat Vector for uniform result comparison.
6. **Result Comparison:** Validation of output correctness across serial, parallel, and PETSc implementations to ensure computational integrity.
7. **Finalisation:** Termination of MPI and PETSc environments and release of resources.

Comprehensive testing, facilitated by the `batch_test.sh` script (see appendix B.H.2), was conducted to evaluate algorithmic performance under varying conditions—spanning different sparse matrix characteristics, fat vector column numbers, and process counts. These tests were categorised into three batches to pinpoint performance influencers:

- **Matrix Impact:** Examining how variations in sparse matrix properties affect algorithm efficiency.
- **Fat Vector Influence:** Assessing the performance implications of altering fat vector column quantities.
- **Execution Time Measurement:** Repeating the second batch tests without tracking average communication and computation times to capture precise execution durations.

This structured approach enables a thorough understanding of each algorithm's behavior under diverse computational scenarios and establishes a foundation for optimising sparse matrix-vector multiplication operations.

# Chapter 3

## Results and Discussion

### 3.1 Results

#### 3.1.1 Sparse Matrix Impact

The first set of experiments focused on the impact of the sparse matrix on the performance of the algorithms. The number of columns in the fat vector was fixed at 6. The performance of all algorithms was evaluated using the following five sparse matrix (2, 3):

Matrix Name	Dimensions	Non-Zero Elements	Symmetric	Type
Cage4	$9 \times 9$	49	No	Real
FEM_3D_thermal1	$17,880 \times 17,880$	430,740	No	Real
DC1	$116,835 \times 116,835$	766,396	No	Real
Cop20k_A	$121,192 \times 121,192$	2,624,331	Yes	Real
Amazon0302	$262,111 \times 262,111$	1,234,877	No	Binary

Table 3.1: Sparse matrix specifications



### 3.1.1.1 Execution Time Evolution

The following figures illustrate the evolution of the execution time for each of the five test matrices as the number of processes increases.




Figure 3.1: Cage4 Matrix Execution Time Evolution




Figure 3.2: DC1 Matrix Execution Time Evolution




Figure 3.3: Amazon0302 Matrix Execution Time Evolution




Figure 3.4: Cop20k\_A Matrix Execution Time Evolution




Figure 3.5: FEM\_3D\_thermal1 Matrix Execution Time Evolution

### 3.1.1.2 Average Communication Time Evolution

The following figures illustrate the evolution of the average communication time per process for each of the five test matrices as the number of processes increases.




Figure 3.6: Cage4 Matrix Communication Time Evolution

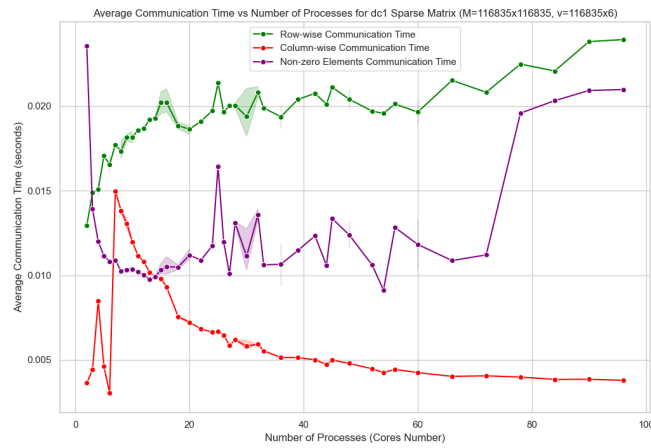


Figure 3.7: DC1 Matrix Communication Time Evolution

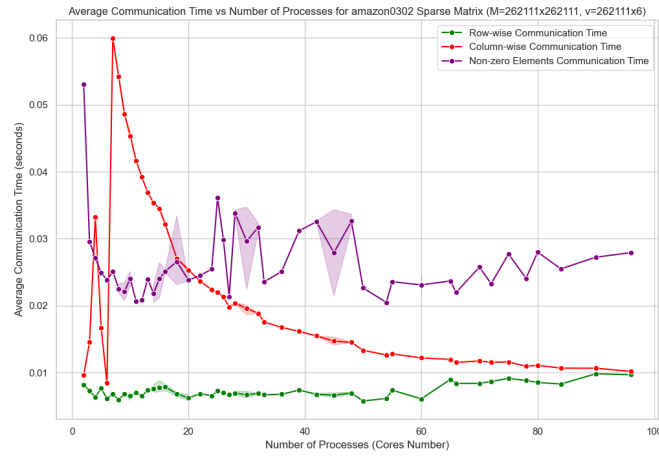


Figure 3.8: Amazon0302 Matrix Communication Time Evolution




Figure 3.9: Cop20k\_A Matrix Communication Time Evolution




Figure 3.10: FEM\_3D\_thermal1 Matrix Communication Time Evolution

### 3.1.1.3 Average Computation Time Evolution

The following figures illustrate the evolution of the average computation time per process for each of the five test matrices as the number of processes increases.




Figure 3.11: Cage4 Matrix Computation Time Evolution




Figure 3.12: DC1 Matrix Computation Time Evolution




Figure 3.13: Amazon0302 Matrix Computation Time Evolution




Figure 3.14: Cop20k\_A Matrix Computation Time Evolution

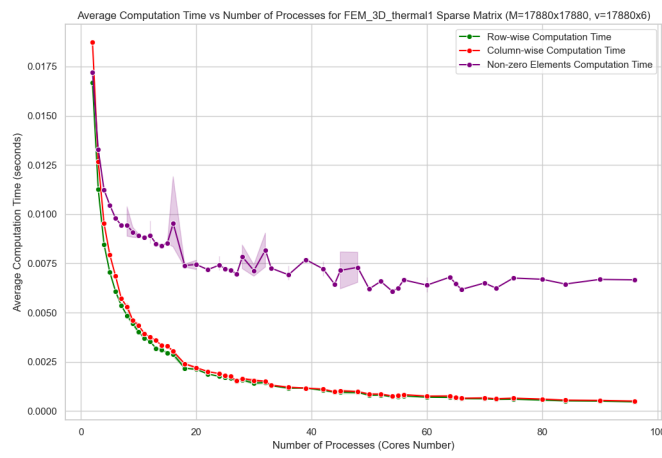


Figure 3.15: FEM\_3D\_thermal1 Matrix Computation Time Evolution

### 3.1.1.4 Performance Evolution

The following figures illustrate the evolution of the performance for each of the five test matrices as the number of processes increases.




Figure 3.16: Cage4 Matrix Performance Evolution




Figure 3.17: DC1 Matrix Performance Evolution




Figure 3.18: Amazon0302 Matrix Performance Evolution




Figure 3.19: Cop20k\_A Matrix Performance Evolution




Figure 3.20: FEM\_3D\_thermal1 Matrix Performance Evolution



### 3.1.1.5 Performance Summary

Here is a summary of the performance of the algorithms:

1. **Serial Algorithm:** The execution time increases with the size and non-zero elements number of the matrix due to the linear increase in computation. As expected, it remains constant with the number of processes. As shown in the figures 3.1, 3.17 and 3.16, the serial algorithm outperforms the parallel algorithms for small matrices or with a lot of non-zero elements non-uniformly distributed like the Cage4 and DC1 matrices. This is due to the overhead associated with coordinating parallel tasks and the communication between processes.
2. **Line-Based Algorithm:** This algorithm performs better with matrices having a balanced distribution of non-zero elements across rows. Matrices with irregular distributions can lead to load imbalance among processes, affecting performance negatively.  
  
For matrices like Amazon0302 or Cop20k\_A with a high number of rows, this algorithm performs well due to its ability to parallelize over rows, but struggles with load balancing for matrices like Cage4 that are smaller or less uniformly distributed. This is evident in the figures 3.16, 3.38 and 3.18. As shown in the figures 3.28 and 3.33, the communication and computation time per process decreases with the number of processes. However, after about 60 processes, the total execution time starts to increase due to the increased communication overhead.
3. **Column-Based Algorithm:** Similar to the row-wise algorithm, but the performance is more sensitive to the matrix's column distribution. Sparsity patterns that lead to dense columns can result in significant overhead due to increased communication between processes.
4. **Non-Zero Element Algorithm:** The performance heavily depends on the distribution of non-zero elements. For matrices with a large number of non-zero elements scattered across, this algorithm can optimise the use of computational resources by focusing work where it's needed. As shown in the figures 3.38 and 3.16, it is efficient for matrices with a high density of non-zero elements like Cop20k\_A, since it focuses computation only where necessary.
5. **PETSc Algorithm:** Based on the execution time, the PETSc algorithm outperforms all the other algorithms (see 3.3). However, the total performance remains constant with the number of processes and very low compared to the other algorithms (see 3.18) as this library is not designed to reconstruct the final result matrix.

### 3.1.2 Fat Vector Impact

The second and third sets of experiments focused on the impact of the dimensions of the fat vector on the performance of the algorithms. The sparse matrix used was Cop20k\_A. The performance of all algorithms was evaluated using the following five fat vector column counts: 1, 3, 6, 9, and 12.

#### 3.1.2.1 Execution Time Evolution

The following figures illustrate the evolution of the execution time for each of the five test fat vector column counts as the number of processes increases.




Figure 3.21: Execution Time Evolution (k=1)




Figure 3.22: Execution Time Evolution (k=3)




Figure 3.23: Execution Time Evolution (k=6)




Figure 3.24: Execution Time Evolution (k=9)




Figure 3.25: Execution Time Evolution (k=12)

### 3.1.2.2 Average Communication Time Evolution

The following figures illustrate the evolution of the average communication time per process for each of the five test fat vector column counts as the number of processes increases.




Figure 3.26: Communication Time Evolution (k=1)




Figure 3.27: Communication Time Evolution (k=3)

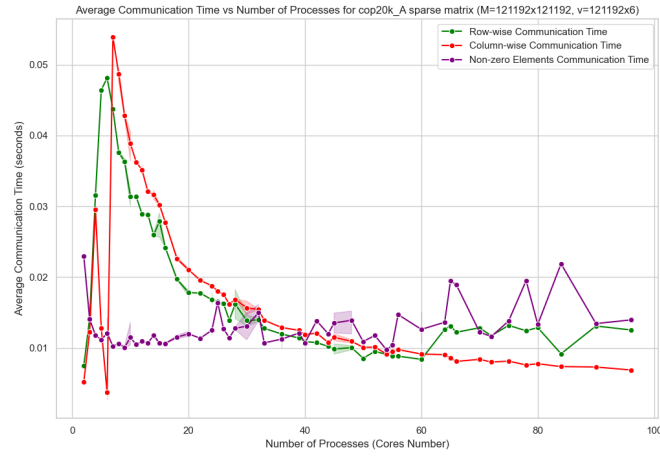


Figure 3.28: Communication Time Evolution (k=6)




Figure 3.29: Communication Time Evolution (k=9)




Figure 3.30: Communication Time Evolution (k=12)

### 3.1.2.3 Average Computation Time Evolution

The following figures illustrate the evolution of the average computation time per process for each of the five test fat vector column counts as the number of processes increases.




Figure 3.31: Computation Time Evolution (k=1)

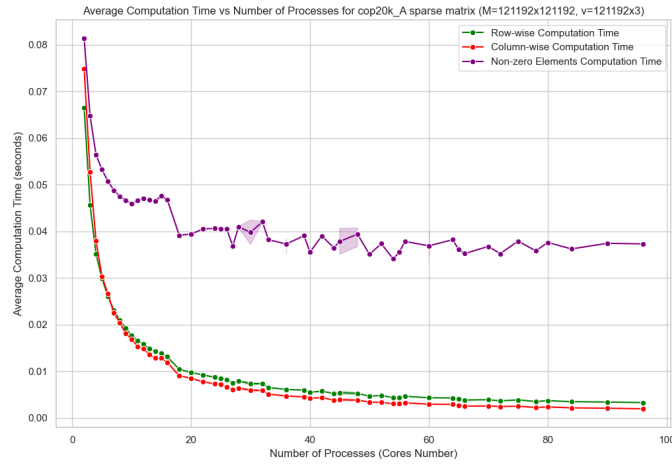


Figure 3.32: Computation Time Evolution (k=3)




Figure 3.33: Computation Time Evolution (k=6)




Figure 3.34: Computation Time Evolution (k=9)




Figure 3.35: Computation Time Evolution (k=12)

### 3.1.2.4 Performance Evolution

The following figures illustrate the evolution of the performance for each of the five test fat vector column counts as the number of processes increases.




Figure 3.36: Performance Evolution (k=1)




Figure 3.37: Performance Evolution (k=3)






Figure 3.38: Performance Evolution (k=6)




Figure 3.39: Performance Evolution (k=9)




Figure 3.40: Performance Evolution (k=12)

### 3.1.2.5 Performance Summary

Here is a summary of the performance of the algorithms:

1. **Serial Algorithm:** As shown in figures 3.36 and 3.40, the serial algorithm outperforms the parallel algorithms for small dimensions but is outperformed by the parallel algorithms for larger dimensions.
2. **Line-Based Algorithm:** As shown in the figures 3.36 and 3.40, the performance of the line-based algorithm increases with the number of processes and the number of columns in the fat vector. However, there is a drop in performance when the number of processes is too large. This is due to the increased communication overhead as shown in the figures 3.26 and 3.30.
3. **Column-Based Algorithm:** As shown in the figures 3.36 and 3.40, the performance of the column-based algorithm increases with the number of processes and the number of columns in the fat vector. However, there is a large drop in performance when the number of processes is greater than the number of columns in the fat vector. This is due to the increased communication overhead as shown in the figures 3.21 and 3.25.
4. **Non-Zero Element Algorithm:** As shown in the figures 3.27 and 3.30, the communication time per process increase proportionally with the number of processes and the number of columns in the fat vector. As the complexity of the computation increases with the number of columns in the fat vector, the performance becomes the same as the row-wise algorithm. This is evident in the figures 3.37 and 3.40.
5. **PETSc Algorithm:** As expected, the execution time decreases with the number of processes but stay constant with the number of columns in the fat vector as shown in the figures 3.21 and 3.25. This is due to the highly optimized parallel sparse matrix-vector multiplication implementation in the PETSc library. However, the performance is not good compared to the custom implementations as it not optimised to reconstruct the final result from the gathered vector.

## 3.2 HPC's Environmental Impact

Although HPC systems are essential for complex calculations and simulations in a variety of fields, their impact on the environment is not negligible. Energy consumption and heat generation require the use of large cooling systems, which further increases their ecological footprint. However, advances in energy-efficient cooling technologies and methods are paving the way for more sustainable HPC operations.

Irina Kupiainen's article highlights the environmental considerations involved in HPC operations, highlighting the exemplary case of LUMI in integrating sustainability. LUMI is one of the most powerful supercomputers in the world, while having a neutral, or even positive, impact on the environment. Located in Finland, it benefits from the country's cool climate and access to renewable energy sources, making it possible to reconcile high performance with environmental sustainability (4).

New innovations are emerging in the field of green HPC. For example, it is possible to use new cooling systems that significantly lower junction temperatures and require less pumping power (5). In addition, Green Cloud Computing initiatives aim to optimise the use of resources, by exploiting renewable energy sources and adopting energy-saving technologies such as virtualisation (6, 7). These advances demonstrate a commitment to sustainability, underlining the move towards reducing energy consumption and emissions in high-performance computing operations, while maintaining energy efficiency.

Finally, the adoption of metrics and standards plays a pivotal role in guiding HPC towards sustainability. The Green Index (TGI) exemplifies this by providing a comprehensive metric for evaluating the energy efficiency of HPC systems (8). It aggregates various benchmarks into a singular measure, facilitating comparisons of system-wide efficiency. Additionally, the Green500 list ranks supercomputers based on their energy efficiency, promoting an industry standard that encourages the development of environmentally friendly HPC technologies. These metrics and rankings not only spotlight the most energy-efficient systems but also drive competition and innovation towards greener HPC solutions.

# Chapter 4

## Conclusion

In conclusion, this report has meticulously analysed various parallelization strategies for fat matrix vector multiplication, based on the MPI and PETSc frameworks, and through extensive experimentation and comparison, highlights the importance of choosing an appropriate parallelization approach based on matrix characteristics and computational resources. The results highlight the trade-offs between computational and communication overheads, emphasising the need for a balanced distribution of workloads to maximise efficiency. As expected the PETSc library outperforms the custom implementations in terms of execution time but not in terms of performance, as the matrices stay distributed even after the multiplication. The custom implementations are more efficient to reconstruct the final result from the gathered vector. In addition, the discussion of the environmental impact of HPC practices, with a nod to sustainable computing, reflects a broader perspective on the implications of computational research. Future directions could explore adaptive parallelization techniques, further optimizing specific models and matrix sizes, possibly integrating AI for dynamic algorithm selection based on real-time performance metrics. This study not only contributes to the field of HPC by providing insight into efficient algorithmic implementations, but also paves the way for more energy-efficient HPC systems, in line with global sustainable development goals.

# References

1. Balay S, Abhyankar S, Adams MF, Benson S, Brown J, Brune P, et al.. PETSc Web page; 2023. <https://petsc.org/>. Available at: <https://petsc.org/>. (Accessed: December 28, 2023).
2. Lugowski A. fast matrix market: Fast and Full-Featured Matrix Market I/O Library; 2023. Available at: [https://github.com/alugowski/fast\\_matrix\\_market](https://github.com/alugowski/fast_matrix_market). (Accessed: December 28, 2023).
3. Kolodziej SP, Aznaveh M, Bullock M, David J, Davis TA, Henderson M, et al. The SuiteSparse Matrix Collection Website Interface. *Journal of Open Source Software*. 2019;4(35):1244-8. Available at: <https://sparse.tamu.edu/>. (Accessed: December 28, 2023).
4. Kupiainen I. HPC at the core of green and digital transition; 2021. Available at: <https://www.scientific-computing.com/viewpoint/hpc-core-green-and-digital-transition>. (Accessed: January 02, 2024).
5. Karwa N. Ultra-Low Global Warming Potential Heat Transfer Fluids for Pumped Two-Phase Cooling in HPC Data Centers; 2020. Available at: <https://ieeexplore.ieee.org/document/9190269>. (Accessed: December 28, 2023).
6. Geetanjali, Quraishi SJ. Energy Savings using Green Cloud Computing; 2022. Available at: <https://ieeexplore.ieee.org/document/9917654>. (Accessed: December 28, 2023).
7. Guyon D, Orgerie AC, Morin C, Agarwal D. How Much Energy Can Green HPC Cloud Users Save?; 2017. Available at: <https://ieeexplore.ieee.org/document/7912681>. (Accessed: December 28, 2023).
8. Subramaniam B, Feng Wc. The Green Index: A Metric for Evaluating System-Wide Energy Efficiency in HPC Systems; 2012. Available at: <https://ieeexplore.ieee.org/document/6270748>. (Accessed: December 28, 2023).

# Appendix A

## Documentation

### Appendix A.A Project tree

```
Source Code/
  scripts /
    batch_test.sh
    get_csv_all.sh
    get_csv_debug.sh
    get_csv_specific.sh
    mpi.sub
  MatrixDefinitions.h
  SparseMatrixFatVectorMultiply.h
  SparseMatrixFatVectorMultiply.cpp
  SparseMatrixFatVectorMultiplyRowWise.h
  SparseMatrixFatVectorMultiplyRowWise.cpp
  SparseMatrixFatVectorMultiplyColumnWise.h
  SparseMatrixFatVectorMultiplyColumnWise.cpp
  SparseMatrixFatVectorMultiplyNonZeroElement.h
  SparseMatrixFatVectorMultiplyNonZeroElement.cpp
  utils.h
  utils.cpp
  main.cpp
results /
  fat_vector_dim /
    <sparse_matrix>_<k>_<metric>.png
  matrix_dim /
    <sparse_matrix>_<k>_<metric>.png
```

### Appendix A.B Getting Started

To run the program, follow these steps:

1. Install the required libraries: mpi & petsc
2. Compile the main program using the following command:

```

        mmpicxx -o <executable_name> -I${PETSC_DIR}/include -I${
PETSC_DIR}/${PETSC_ARCH}/include -L${PETSC_DIR}/${PETSC_ARCH}/lib -lpetsc
SparseMatrixFatVectorMultiply.cpp main_verify.cpp utils.cpp
SparseMatrixFatVectorMultiplyColumnWise.cpp
SparseMatrixFatVectorMultiplyNonZeroElement.cpp
SparseMatrixFatVectorMultiplyRowWise.cpp

```

3. Run the program using the following command:

```

        mpirun -np <number_of_processes> <executable_name> <k> <
sparse_matrix_file_pathw>

```

## Appendix A.C Methods Overview

### A.C.1 Utils.h

#### A.C.1.1 ConvertPETScMatToFatVector

**Description:** Converts a PETSc matrix to a FatVector stucture.

**Parameters:**

- Mat C - PETSc matrix to be converted.

**Returns:** FatVector - Fat vector representation of the PETSc matrix.

#### A.C.1.2 areMatricesEqual

**Description:** Compares two matrices for equality within a specified tolerance.

**Parameters:**

- FatVector &mat1 - First matrix.
- FatVector &mat2 - Second matrix.
- double tolerance - Tolerance for comparison.

**Returns:** bool - True if matrices are equal within the tolerance, false otherwise.

#### A.C.1.3 readMatrixMarketFile

**Description:** Reads a matrix from a Matrix Market file into a sparse matrix format.

**Parameters:**

- std::string &filename - Name of the Matrix Market file.

**Returns:** SparseMatrix - Sparse matrix read from the file.

#### A.C.1.4 generateLargeFatVector

**Description:** Generates a random Fat Vector with specified dimensions.

**Parameters:**

- `int n` - Number of rows.
- `int k` - Number of columns.

**Returns:** `FatVector` - Generated fat vector.

#### A.C.1.5 serialize and deserialize

**Description:** Serializes and deserializes a `FatVector` to and from a flat array, respectively.

**Parameters for serialize:**

- `FatVector &fatVec` - fat vector to serialize.

**Returns:** `std::vector<double>` - Flat array containing the serialized data.

**Parameters for deserialize:**

- `std::vector<double> &flat` - Flat array to deserialize.
- `int rows` - Number of rows in the fat vector.
- `int cols` - Number of columns in the fat vector.

**Returns:** `FatVector` - Deserialized fat vector.

### A.C.2 SparseMatrixFatVectorMultiply.h

#### A.C.2.1 sparseMatrixFatVectorMultiply

**Description:** Executes the multiplication using a sequential algorithm.

**Parameters:**

- `SparseMatrix &sparseMatrix` - The sparse matrix.
- `FatVector &fatVector` - The Fat Vector.
- `int vecCols` - Number of columns in the Fat Vector.

**Returns:** `FatVector` - Result of the multiplication.



### A.C.3 SparseMatrixFatVectorMultiplyRowWise.h

#### A.C.3.1 sparseMatrixFatVectorMultiplyRowWise

**Description:** Multiplies a sparse matrix with a Fat Vector using row-wise distribution.

**Parameters:**

- SparseMatrix &sparseMatrix - The sparse matrix.
- FatVector &fatVector - The Fat Vector.
- int vecCols - Number of columns in the Fat Vector.

**Returns:** FatVector - Result of the multiplication.

### A.C.4 SparseMatrixFatVectorMultiplyColumnWise.h

#### A.C.4.1 sparseMatrixFatVectorMultiplyColumnWise

**Description:** Executes the multiplication using column-wise parallel algorithm.

**Parameters:**

- SparseMatrix &sparseMatrix - The sparse matrix.
- FatVector &fatVector - The Fat Vector.
- int vecCols - Number of columns in the Fat Vector.

**Returns:** FatVector - Result of the multiplication.

### A.C.5 SparseMatrixFatVectorMultiplyNonZeroElement.h

#### A.C.5.1 sparseMatrixFatVectorMultiplyNonZeroElement

**Description:** Executes the multiplication using non-zero element parallel algorithm.

**Parameters:**

- SparseMatrix &sparseMatrix - The sparse matrix.
- FatVector &fatVector - The Fat Vector.
- int vecCols - Number of columns in the Fat Vector.

**Returns:** FatVector - Result of the multiplication.

# Appendix B

## Source Codes

### Appendix B.A Data Structures

Data stuctures of the sparse matrix and fat vector.

```
1  #ifndef MATRIXDEFINITIONS_H
2  #define MATRIXDEFINITIONS_H
3
4  #include <vector>
5
6  //
7  /**
8   * @brief Struct to represent a sparse matrix
9   *
10  * @param values    Non-zero values
11  * @param colIndices Column indices of non-zero values
12  * @param rowPtr    Row pointers
13  */
14  struct SparseMatrix
15  {
16      std::vector<double> values;
17      std::vector<int> colIndices;
18      std::vector<int> rowPtr;
19  };
20
21  // Type definition for a Fat Vector
22  typedef std::vector<std::vector<double>> FatVector;
23
24  #endif
```

## Appendix B.B Sequential Algorithm

Sequential algorithm for multiplying a sparse matrix by a fat vector.

### B.B.1 Declaration File

```

1 #ifndef SPARSEMATRIXFATVECTORMULTIPLY_H
2 #define SPARSEMATRIXFATVECTORMULTIPLY_H
3
4 #include "MatrixDefinitions.h"
5
6 /**
7  * @brief Function to execute the sparse matrix-fat vector multiplication using
8  * sequential algorithm
9  *
10  * @param sparseMatrix Sparse matrix
11  * @param fatVector Fat Vector
12  * @param vecCols Number of columns in the Fat Vector
13  * @return FatVector Result of the multiplication
14  */
15 FatVector sparseMatrixFatVectorMultiply(const SparseMatrix &sparseMatrix,
16                                         const FatVector &fatVector, int vecCols
17                                         );
18 #endif

```

### B.B.2 Implementation File

```

1 #include "SparseMatrixFatVectorMultiply.h"
2
3 /**
4  * @brief Function to execute the sparse matrix-Fat Vector multiplication using
5  * sequential algorithm
6  *
7  * @param sparseMatrix Sparse matrix
8  * @param fatVector Fat Vector
9  * @param vecCols Number of columns in the Fat Vector
10  * @return FatVector Result of the multiplication
11  */
12 FatVector sparseMatrixFatVectorMultiply(const SparseMatrix &sparseMatrix,
13                                         const FatVector &fatVector, int vecCols
14                                         )
15 {
16     // Initialisation of the result vector
17     FatVector result(sparseMatrix.numRows, std::vector<double>(vecCols, 0.0));
18     // Iterate over the rows of the sparse matrix
19     for (int i = 0; i < sparseMatrix.numRows; ++i)
20     {
21         // Iterate over the non-zero elements in the current row
22         for (int j = sparseMatrix.rowPtr[i]; j < sparseMatrix.rowPtr[i + 1]; ++j)
23         {
24             // Iterate over the columns of the Fat Vector
25             for (int k = 0; k < vecCols; ++k)
26             {
27                 result[i][k] += sparseMatrix.values[j] * fatVector[sparseMatrix.colIndices[j]][k]; // Compute the result
28             }
29         }
30     }
31     // Return the result
32     return result;
33 }

```

## Appendix B.C Line-Based Parallelism

Parallel algorithm for multiplying a sparse matrix by a fat vector using line-based parallelism.

### B.C.1 Declaration File

```

1  #ifndef SPARSEMATRIXFATVECTORMULTIPLYROWWISE_H
2  #define SPARSEMATRIXFATVECTORMULTIPLYROWWISE_H
3
4  #include "MatrixDefinitions.h"
5  #include <iostream> // std::cout
6
7  /**
8   * @brief Function to multiply a sparse matrix with a Fat Vector using row-wise
      distribution
9   *
10  * @param sparseMatrix The sparse matrix to be multiplied
11  * @param fatVector The Fat Vector to be multiplied
12  * @param vecCols Number of columns in the Fat Vector
13  * @return FatVector Result of the multiplication
14  */
15  FatVector sparseMatrixFatVectorMultiplyRowWise(const SparseMatrix &sparseMatrix,
16                                                  const FatVector &fatVector,
17                                                  int vecCols);
18
19 #endif

```

### B.C.2 Implementation File

```

1  #include <mpi.h>
2  #include "SparseMatrixFatVectorMultiplyRowWise.h"
3
4  /**
5   * @brief Function to multiply a sparse matrix with a Fat Vector using row-wise
      distribution
6   *
7   * @param sparseMatrix The sparse matrix to be multiplied
8   * @param fatVector The Fat Vector to be multiplied
9   * @param vecCols Number of columns in the Fat Vector
10  * @return FatVector Result of the multiplication
11  */
12  FatVector sparseMatrixFatVectorMultiplyRowWise(const SparseMatrix &sparseMatrix,
13                                                  const FatVector &fatVector,
14                                                  int vecCols)
15  {
16      // Retrieve the rank and size of the MPI world
17      int worldSize, worldRank;
18      MPI_Comm_size(MPI_COMM_WORLD, &worldSize);
19      MPI_Comm_rank(MPI_COMM_WORLD, &worldRank);
20
21      // ===== FOR DEBUGGING ONLY - START LOCAL COMPUTATION
22      // double computation_start = MPI_Wtime();
23      // ===== FOR DEBUGGING ONLY - START LOCAL COMPUTATION
24
25      // Distribute rows among processes
26      int rowsCountPerProcess = sparseMatrix.numRows / worldSize;
27      // Number of rows per process
28      int extraRows = sparseMatrix.numRows % worldSize;
29      // Number of extra rows to be distributed
30      // among processes

```

```

28     int startRow = worldRank * rowCountPerProcess + std::min(worldRank, extraRows)
29     ; // Starting row index for the current process
30     int endRow = startRow + rowCountPerProcess + (worldRank < extraRows ? 1 : 0);
31     // Ending row index for the current process
32
33     // Local computation
34     int localSize = (endRow - startRow) * vecCols; // Number of elements in the
35     local result vector
36     std::vector<double> localResult(localSize); // Local result vector
37
38     // Iterate over the rows assigned to the current process
39     for (int i = startRow; i < endRow; ++i)
40     {
41         // Iterate over the non-zero elements in the current row
42         for (int j = sparseMatrix.rowPtr[i]; j < sparseMatrix.rowPtr[i + 1]; ++j)
43         {
44             int colIndex = sparseMatrix.colIndices[j]; // Column index of the non-
45             zero element
46
47             // Iterate over the columns of the Fat Vector
48             for (int k = 0; k < vecCols; ++k)
49             {
50                 int localIndex = (i - startRow) * vecCols + k;
51                 // Index of the element in the
52                 local result vector
53                 localResult[localIndex] += sparseMatrix.values[j] * fatVector[
54                 colIndex][k]; // Compute the result
55             }
56         }
57     }
58
59     // ===== FOR DEBUGGING ONLY - STOP LOCAL COMPUTATION
60     TIMER =====
61     // double computation_end = MPI_Wtime();
62     // double local_computation_time = computation_end - computation_start;
63     // ===== FOR DEBUGGING ONLY - STOP LOCAL COMPUTATION
64     TIMER =====
65
66     // ===== FOR DEBUGGING ONLY - START COMMUNICATION TIMER
67     =====
68     // Start timing for communication
69     // double communication_start = MPI_Wtime();
70     // ===== FOR DEBUGGING ONLY - START COMMUNICATION TIMER
71     =====
72
73     // Preparation for Gather operation
74     std::vector<int> recvCounts(worldSize), displacements(worldSize);
75     if (worldRank == 0)
76     {
77         int totalSize = 0; // Total number of elements to be received
78
79         // Compute the number of elements to be received from each process
80         for (int rank = 0; rank < worldSize; ++rank)
81         {
82             int startRowThisRank = rank * rowCountPerProcess + std::min(rank,
83             extraRows); // Starting row index for the current
84             process
85             int endRowThisRank = startRowThisRank + rowCountPerProcess + (rank <
86             extraRows ? 1 : 0); // Ending row index for the current process
87             recvCounts[rank] = (endRowThisRank - startRowThisRank) * vecCols;
88             // Number of elements to be received from
89             the current process
90             displacements[rank] = totalSize;
91             //
92             Displacement for the current process
93             totalSize += recvCounts[rank];
94             //
95             Update the total number of elements to be received
96         }
97     }

```

```

78
79 // Gather all local results into the root process
80 std::vector<double> gatheredResults;
81 if (worldRank == 0)
82 {
83     gatheredResults.resize(recvCounts[0] * worldSize); // Resize the vector to
84     hold all the results
85 }
86 MPI_Gatherv(localResult.data(), localSize, MPI_DOUBLE,
87             gatheredResults.data(), recvCounts.data(),
88             displacements.data(), MPI_DOUBLE, 0, MPI_COMM_WORLD); // Gather the
89             local results in the root process
90
91 // ===== FOR DEBUGGING ONLY - STOP COMMUNICATION TIMER
92 // =====
93 // double communication_end = MPI_Wtime();
94 // double local_communication_time = communication_end - communication_start;
95 // ===== FOR DEBUGGING ONLY - STOP COMMUNICATION TIMER
96 // =====
97
98 // ===== FOR DEBUGGING ONLY - COLLECTING AND ANALYSING
99 // PERFORMANCE DATA =====
100 // double total_computation_time = 0.0, total_communication_time = 0.0;
101 // MPI_Reduce(&local_computation_time, &total_computation_time, 1, MPI_DOUBLE,
102 // MPI_SUM, 0, MPI_COMM_WORLD);
103 // MPI_Reduce(&local_communication_time, &total_communication_time, 1,
104 // MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
105 // ===== FOR DEBUGGING ONLY - COLLECTING AND ANALYSING
106 // PERFORMANCE DATA =====
107
108 // Reconstruct the final result matrix in the root process
109 FatVector finalResult;
110 if (worldRank == 0)
111 {
112     // ===== FOR DEBUGGING ONLY - PRINTING PERFORMANCE
113     // DATA =====
114     // double avg_computation_time = total_computation_time / worldSize;
115     // double avg_communication_time = total_communication_time / worldSize;
116     // std::cout << "Row-wise Average Computation Time: " <<
117     // avg_computation_time << std::endl;
118     // std::cout << "Row-wise Average Communication Time: " <<
119     // avg_communication_time << std::endl;
120     // ===== FOR DEBUGGING ONLY - PRINTING PERFORMANCE
121     // DATA =====
122
123     finalResult.resize(sparseMatrix.numRows, std::vector<double>(vecCols, 0.0))
124     ; // Resize the final result matrix
125
126     // Iterate over the rows of the final result
127     for (int i = 0, index = 0; i < sparseMatrix.numRows; ++i)
128     {
129         // Iterate over the columns of the final result
130         for (int j = 0; j < vecCols; ++j, ++index)
131         {
132             finalResult[i][j] = gatheredResults[index]; // Copy the element of
133             the final result
134         }
135     }
136 }
137
138 // Return the final result
139 return (worldRank == 0) ? finalResult : FatVector{};
140 }

```

## Appendix B.D Column-Wise Parallelism

Parallel algorithm for multiplying a sparse matrix by a fat vector using column-wise parallelism.

### B.D.1 Declaration File

```

1 #ifndef SPARSEMATRIXFATVECTORMULTIPLYCOLUMNWISE_H
2 #define SPARSEMATRIXFATVECTORMULTIPLYCOLUMNWISE_H
3
4 #include "MatrixDefinitions.h"
5 #include <iostream> // std::cout
6
7 /**
8  * @brief Function to execute the sparse matrix-Fat Vector multiplication using
9  * column-wise parallel algorithm
10  *
11  * @param sparseMatrix Sparse matrix
12  * @param fatVector Fat Vector
13  * @param vecCols Number of columns in the Fat Vector
14  * @return FatVector Result of the multiplication
15  */
16 FatVector sparseMatrixFatVectorMultiplyColumnWise(const SparseMatrix &sparseMatrix,
17     const FatVector &fatVector, int vecCols);
18 #endif

```

### B.D.2 Implementation File

```

1 #include <mpi.h>
2 #include "SparseMatrixFatVectorMultiplyColumnWise.h"
3 #include <numeric> // std::accumulate
4
5 /**
6  * @brief Function to execute the sparse matrix-Fat Vector multiplication using
7  * column-wise parallel algorithm
8  *
9  * @param sparseMatrix Sparse matrix
10  * @param fatVector Fat Vector
11  * @param vecCols Number of columns in the Fat Vector
12  * @return FatVector Result of the multiplication
13  */
14 FatVector sparseMatrixFatVectorMultiplyColumnWise(const SparseMatrix &sparseMatrix,
15     const FatVector &fatVector, int vecCols)
16 {
17     // Retrieve the rank and size of the MPI world
18     int worldSize, worldRank;
19     MPI_Comm_size(MPI_COMM_WORLD, &worldSize);
20     MPI_Comm_rank(MPI_COMM_WORLD, &worldRank);
21
22     // ===== FOR DEBUGGING ONLY - START LOCAL COMPUTATION
23     // double computation_start = MPI_Wtime();
24     // ===== FOR DEBUGGING ONLY - START LOCAL COMPUTATION
25
26     // Distribute columns among processes
27     int colsPerProcess = vecCols / worldSize;
28
29     // Number of columns per process
30     int extraCols = vecCols % worldSize;
31
32     // Number of extra columns to be distributed among processes
33 }

```

```

27     int startCol = worldRank * colsPerProcess;
28                                     //
29     Starting column index for the current process
30     int endCol = (worldRank != worldSize - 1) ? startCol + colsPerProcess :
31     startCol + colsPerProcess + extraCols; // Ending column index for the
32     current process
33
34     // Local computation
35     int localSize = sparseMatrix.numRows * (endCol - startCol); // Number of
36     elements in the local result vector
37     std::vector<double> localResult(localSize, 0.0);
38     // Iterate over the columns assigned to the current process
39     for (int col = startCol; col < endCol; ++col)
40     {
41         // Iterate over the rows of the sparse matrix
42         for (int i = 0; i < sparseMatrix.numRows; ++i)
43         {
44             // Iterate over the non-zero elements in the current row
45             double sum = 0.0;
46             for (int j = sparseMatrix.rowPtr[i]; j < sparseMatrix.rowPtr[i + 1]; ++j)
47             {
48                 int sparseCol = sparseMatrix.colIndices[j]; //
49                 Column index of the non-zero element
50                 sum += sparseMatrix.values[j] * fatVector[sparseCol][col]; //
51                 Compute the result
52             }
53             localResult[i * (endCol - startCol) + (col - startCol)] = sum; // Store
54             the result in the local result vector
55         }
56     }
57
58     // ===== FOR DEBUGGING ONLY - STOP LOCAL COMPUTATION
59     TIMER =====
60     // double computation_end = MPI_Wtime();
61     // double local_computation_time = computation_end - computation_start;
62     // ===== FOR DEBUGGING ONLY - STOP LOCAL COMPUTATION
63     TIMER =====
64
65     // ===== FOR DEBUGGING ONLY - START COMMUNICATION TIMER
66     =====
67     // Start timing for communication
68     // double communication_start = MPI_Wtime();
69     // ===== FOR DEBUGGING ONLY - START COMMUNICATION TIMER
70     =====
71
72     // Preparation for Gather operation
73     std::vector<int> recvCounts(worldSize), displacements(worldSize); // Number of
74     elements to be received from each process, Displacement for each process
75     if (worldRank == 0)
76     {
77         // Compute the number of elements to be received from each process
78         int displacement = 0;
79         for (int i = 0; i < worldSize; ++i)
80         {
81             int startColThisRank = i * colsPerProcess;
82
83             // Starting column index for the current process
84             int endColThisRank = (i != worldSize - 1) ? startColThisRank +
85             colsPerProcess : startColThisRank + colsPerProcess + extraCols; //
86             Ending column index for the current process
87             recvCounts[i] = sparseMatrix.numRows * (endColThisRank -
88             startColThisRank);
89                                     // Number of
90             elements to be received from the current process
91             displacements[i] = displacement;
92
93             // Displacement for the current process
94             displacement += recvCounts[i];

```



```

73         // Update the displacement
74     }
75 }
76 // Gather all local results into the root process
77 std::vector<double> gatheredResults;
78 if (worldRank == 0)
79 {
80     gatheredResults.resize(std::accumulate(recvCounts.begin(), recvCounts.end(),
81     , 0)); // Resize the vector to hold the final result
82 }
83 MPI_Gatherv(localResult.data(), localSize, MPI_DOUBLE,
84     gatheredResults.data(), recvCounts.data(),
85     displacements.data(), MPI_DOUBLE, 0, MPI_COMM_WORLD); // Gather the
86     local results in the root process
87
88 // ===== FOR DEBUGGING ONLY - STOP COMMUNICATION TIMER
89 // =====
90 // double communication_end = MPI_Wtime();
91 // double local_communication_time = communication_end - communication_start;
92 // ===== FOR DEBUGGING ONLY - STOP COMMUNICATION TIMER
93 // =====
94
95 // ===== FOR DEBUGGING ONLY - COLLECTING AND ANALYSING
96 // PERFORMANCE DATA =====
97 // double total_computation_time = 0.0, total_communication_time = 0.0;
98 // MPI_Reduce(&local_computation_time, &total_computation_time, 1, MPI_DOUBLE,
99 // MPI_SUM, 0, MPI_COMM_WORLD);
100 // MPI_Reduce(&local_communication_time, &total_communication_time, 1,
101 // MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
102 // ===== FOR DEBUGGING ONLY - COLLECTING AND ANALYSING
103 // PERFORMANCE DATA =====
104
105 // Reconstruct the final result matrix in the root process
106 FatVector finalResult;
107 if (worldRank == 0)
108 {
109     // ===== FOR DEBUGGING ONLY - PRINTING PERFORMANCE
110     // DATA =====
111     // double avg_computation_time = total_computation_time / worldSize;
112     // double avg_communication_time = total_communication_time / worldSize;
113     // std::cout << "Column-wise Average Computation Time: " <<
114     //     avg_computation_time << std::endl;
115     // std::cout << "Column-wise Average Communication Time: " <<
116     //     avg_communication_time << std::endl;
117     // ===== FOR DEBUGGING ONLY - PRINTING PERFORMANCE
118     // DATA =====
119
120     // Reconstruct the final result matrix
121     finalResult.resize(sparseMatrix.numRows, std::vector<double>(vecCols, 0.0))
122     ; // Resize the final result matrix
123     int resultIndex = 0;
124     // Iterate over the processes
125     for (int rank = 0; rank < worldSize; ++rank)
126     {
127         int numColsThisRank = (rank != worldSize - 1) ? colsPerProcess :
128             colsPerProcess + extraCols; // Number of columns assigned to the
129             current process
130         int startColThisRank = rank * colsPerProcess;
131
132         // Starting column
133         // index for the current process
134
135         // Iterate over the rows of the sparse matrix
136         for (int row = 0; row < sparseMatrix.numRows; ++row)
137         {
138             // Iterate over the columns assigned to the current process
139             for (int col = 0; col < numColsThisRank; ++col)
140             {
141                 finalResult[row][startColThisRank + col] = gatheredResults[
142                     resultIndex++]; // Reconstruct the final result matrix

```

```

124         }
125     }
126 }
127 }
128
129 // Return the final result
130 return (worldRank == 0) ? finalResult : FatVector{};
131 }

```

## Appendix B.E Non-Zero Element Parallelism

Parallel algorithm for multiplying a sparse matrix by a fat vector using non-zero element parallelism.

### B.E.1 Declaration File

```

1  #ifndef SPARSEMATRIXFATVECTORMULTIPLYNONZEROELEMENT_H
2  #define SPARSEMATRIXFATVECTORMULTIPLYNONZEROELEMENT_H
3
4  #include "MatrixDefinitions.h"
5  #include <iostream> // std::cout
6
7  /**
8   * @brief Function to execute the sparse matrix-Fat Vector multiplication using non
9   *         -zero element parallel algorithm
10  *
11  * @param sparseMatrix Sparse matrix
12  * @param fatVector Fat Vector
13  * @param vecCols Number of columns in the Fat Vector
14  * @return FatVector Result of the multiplication
15  */
16 FatVector sparseMatrixFatVectorMultiplyNonZeroElement(const SparseMatrix &
17     sparseMatrix, const FatVector &fatVector, int vecCols);
18 #endif

```

### B.E.2 Implementation File

```

1  #include "SparseMatrixFatVectorMultiplyNonZeroElement.h"
2  #include <mpi.h>
3
4  /**
5   * @brief Function to execute the sparse matrix-Fat Vector multiplication using non
6   *         -zero element parallel algorithm
7   *
8   * @param sparseMatrix Sparse matrix
9   * @param fatVector Fat Vector
10  * @param vecCols Number of columns in the Fat Vector
11  * @return FatVector Result of the multiplication
12  */
13 FatVector sparseMatrixFatVectorMultiplyNonZeroElement(const SparseMatrix &
14     sparseMatrix, const FatVector &fatVector, int vecCols)
15 {
16     // Retrieve the rank and size of the MPI world
17     int worldSize, worldRank;
18     MPI_Comm_size(MPI_COMM_WORLD, &worldSize);
19     MPI_Comm_rank(MPI_COMM_WORLD, &worldRank);
20
21     // ===== FOR DEBUGGING ONLY - START LOCAL COMPUTATION
22     TIMER =====

```

```

20 // double computation_start = MPI_Wtime();
21 // ===== FOR DEBUGGING ONLY - START LOCAL COMPUTATION
    TIMER =====
22
23 // Distribute non-zero elements among processes
24 int totalNonZeroElements = sparseMatrix.values.size(); // Total number of
    non-zero elements
25 int elementsPerProcess = totalNonZeroElements / worldSize; // Number of non-
    zero elements per process
26 int extraElements = totalNonZeroElements % worldSize; // Number of extra
    non-zero elements to be distributed among processes
27 int startIdx, endIdx; // Starting and
    ending indices of the non-zero elements for the current process
28
29 // Determine the starting and ending indices of the non-zero elements for the
    current process
30 if (worldRank < extraElements)
31 {
32     startIdx = worldRank * (elementsPerProcess + 1); // Add 1 to account for
        the extra non-zero elements
33     endIdx = startIdx + elementsPerProcess + 1; // Add 1 to account for
        the extra non-zero elements
34 }
35 else
36 {
37     startIdx = worldRank * elementsPerProcess + extraElements; // Add
        extraElements to account for the extra non-zero elements
38     endIdx = startIdx + elementsPerProcess; // Add
        extraElements to account for the extra non-zero elements
39 }
40
41 // Map the indices of the non-zero elements to their corresponding row indices
42 std::vector<int> rowIndexMap(sparseMatrix.values.size());
43 // Iterate over the rows of the sparse matrix
44 for (int row = 0, idx = 0; row < sparseMatrix.rowPtr.size() - 1; ++row)
45 {
46     // Iterate over the non-zero elements in the current row
47     for (; idx < sparseMatrix.rowPtr[row + 1]; ++idx)
48     {
49         rowIndexMap[idx] = row; // Map the index of the non-zero element to its
            corresponding row index
50     }
51 }
52
53 // Local computation
54 std::vector<double> localResult(sparseMatrix.numRows * vecCols, 0.0);
55 // Iterate over the non-zero elements assigned to the current process
56 for (int idx = startIdx; idx < endIdx; ++idx)
57 {
58     int row = rowIndexMap[idx]; // Row index of the non-zero
        element
59     int col = sparseMatrix.colIndices[idx]; // Column index of the non-zero
        element
60     double value = sparseMatrix.values[idx]; // Value of the non-zero element
61
62     // Iterate over the columns of the Fat Vector
63     for (int k = 0; k < vecCols; ++k)
64     {
65         localResult[row * vecCols + k] += value * fatVector[col][k]; // Compute
            the result
66     }
67 }
68
69 // ===== FOR DEBUGGING ONLY - STOP LOCAL COMPUTATION
    TIMER =====
70 // double computation_end = MPI_Wtime();
71 // double local_computation_time = computation_end - computation_start;
72 // ===== FOR DEBUGGING ONLY - STOP LOCAL COMPUTATION
    TIMER =====
73

```

```

74 // ===== FOR DEBUGGING ONLY - START COMMUNICATION TIMER
75 // =====
76 // FOR DEBUGGING ONLY - START COMMUNICATION TIMER
77 // double communication_start = MPI_Wtime();
78 // ===== FOR DEBUGGING ONLY - START COMMUNICATION TIMER
79 // =====
80 // Initialise the final result only in the root process
81 FatVector finalResult;
82 if (worldRank == 0)
83 {
84     finalResult.resize(sparseMatrix.numRows, std::vector<double>(vecCols, 0.0))
85     ;
86 }
87 // Gather the local results in the root process
88 std::vector<double> flatFinalResult(sparseMatrix.numRows * vecCols, 0.0);
89 // Flat vector to
90 hold the final result
91 MPI_Reduce(localResult.data(), flatFinalResult.data(), sparseMatrix.numRows *
92 vecCols, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD); // Gather the local
93 results in the root process
94
95 // ===== FOR DEBUGGING ONLY - STOP COMMUNICATION TIMER
96 // =====
97 // double communication_end = MPI_Wtime();
98 // double local_communication_time = communication_end - communication_start;
99 // ===== FOR DEBUGGING ONLY - STOP COMMUNICATION TIMER
100 // =====
101 // ===== FOR DEBUGGING ONLY - COLLECTING AND ANALYSING
102 // PERFORMANCE DATA =====
103 // double total_computation_time = 0.0, total_communication_time = 0.0;
104 // MPI_Reduce(&local_computation_time, &total_computation_time, 1, MPI_DOUBLE,
105 // MPI_SUM, 0, MPI_COMM_WORLD);
106 // MPI_Reduce(&local_communication_time, &total_communication_time, 1,
107 // MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
108 // ===== FOR DEBUGGING ONLY - COLLECTING AND ANALYSING
109 // PERFORMANCE DATA =====
110
111 // Reconstruct the finalResult from flatFinalResult in the root process
112 if (worldRank == 0)
113 {
114     // ===== FOR DEBUGGING ONLY - PRINTING PERFORMANCE
115     // DATA =====
116     // double avg_computation_time = total_computation_time / worldSize;
117     // double avg_communication_time = total_communication_time / worldSize;
118     // std::cout << "Non-zero elements Average Computation Time: " <<
119     // avg_computation_time << std::endl;
120     // std::cout << "Non-zero elements Average Communication Time: " <<
121     // avg_communication_time << std::endl;
122     // ===== FOR DEBUGGING ONLY - PRINTING PERFORMANCE
123     // DATA =====
124
125     // Iterate over the rows of the final result
126     for (int i = 0; i < sparseMatrix.numRows; ++i)
127     {
128         std::copy(flatFinalResult.begin() + i * vecCols, flatFinalResult.begin
129             () + (i + 1) * vecCols, finalResult[i].begin()); // Copy the row of
130             the final result
131     }
132 }
133
134 // Return the final result
135 return (worldRank == 0) ? finalResult : FatVector{};
136 }

```

## Appendix B.F Utility Functions

Utility functions used by the main file.

### B.F.1 Declaration File

```

1  #ifndef UTILS_H
2  #define UTILS_H
3
4  #include <iostream> // std::cout
5  #include <vector>   // std::vector
6  #include <cstdlib>  // rand() and srand()
7  #include <ctime>    // time()
8  #include <mpi.h>
9  #include <petsc.h>
10 #include <fstream>   // std::ifstream
11 #include <string>    // std::string
12 #include <sstream>   // std::stringstream
13 #include <utility>   // std::pair
14 #include <algorithm> // std::sort
15 #include <stdexcept> // std::runtime_error
16 #include <cmath>     // std::fabs
17 #include "MatrixDefinitions.h"
18
19 /**
20  * Method to convert a PETSc matrix to a fat vector
21  * @param C PETSc matrix
22  * @return FatVector fat vector
23  */
24 FatVector ConvertPETScMatToFatVector(Mat C);
25
26 /**
27  * Method to compare two matrices
28  * @param mat1 First matrix
29  * @param mat2 Second matrix
30  * @param tolerance Tolerance for comparison
31  * @return bool True if the matrices are equal, false otherwise
32  */
33 bool areMatricesEqual(const FatVector &mat1, const FatVector &mat2, double
    tolerance);
34
35 /**
36  * Method to read a matrix from a Matrix Market file
37  * @param filename Name of the file
38  * @return SparseMatrix Sparse matrix
39  */
40 SparseMatrix readMatrixMarketFile(const std::string &filename);
41
42 /**
43  * Method to generate a random fat vector
44  * @param n Number of rows
45  * @param m Number of columns
46  * @return FatVector fat vector
47  */
48 FatVector generateLargeFatVector(int n, int k);
49
50 /**
51  * @brief Method to serialize a FatVector to a flat array
52  * @param denseVec fat vector to serialize
53  * @return std::vector<double> Flat array containing the serialized data
54  */
55 std::vector<double> serialize(const FatVector &denseVec);
56
57 /**
58  * @brief Method to deserialize a flat array to a FatVector
59  * @param flat Flat array to deserialize
60  * @param rows Number of rows in the fat vector

```

```

61  * @param cols Number of columns in the fat vector
62  * @return FatVector fat vector
63  */
64  FatVector deserialize(const std::vector<double> &flat, int rows, int cols);
65
66  #endif

```

## B.F.2 Implementation File

```

1  #include "utils.h"
2
3  /**
4   * Method to convert a PETSc matrix to a fat vector
5   * @param C PETSc matrix
6   * @return FatVector fat vector
7   */
8  FatVector ConvertPETScMatToFatVector(Mat C)
9  {
10     PetscInt m, n;           // Number of rows and columns in the matrix
11     MatGetSize(C, &m, &n); // Get the number of rows and columns in the matrix
12
13     FatVector denseVec(m, std::vector<double>(n, 0.0)); // fat vector to hold the
14     matrix
15
16     // Iterate over the rows of the matrix
17     for (int i = 0; i < m; ++i)
18     {
19         // Iterate over the columns of the matrix
20         for (int j = 0; j < n; ++j)
21         {
22             PetscScalar value;           // Value of the element
23             MatGetValues(C, i, j, &value); // Get the value of the element
24             denseVec[i][j] = PetscRealPart(value); // Copy the value of the element
25         }
26     }
27
28     // Return the fat vector
29     return denseVec;
30 }
31
32 /**
33  * Method to compare two matrices
34  * @param mat1 First matrix
35  * @param mat2 Second matrix
36  * @param tolerance Tolerance for comparison
37  * @return bool True if the matrices are equal, false otherwise
38  */
39 bool areMatricesEqual(const FatVector &mat1, const FatVector &mat2, double
40 tolerance)
41 {
42     // Check if the matrices have the same dimensions
43     if (mat1.size() != mat2.size())
44         return false;
45
46     // Iterate over the rows of the matrices
47     for (size_t i = 0; i < mat1.size(); ++i)
48     {
49         // Check if the rows have the same dimensions
50         if (mat1[i].size() != mat2[i].size())
51             return false;
52
53         // Iterate over the columns of the matrices
54         for (size_t j = 0; j < mat1[i].size(); ++j)
55         {
56             // Check if the elements are equal
57             if (std::fabs(mat1[i][j] - mat2[i][j]) > tolerance)
58                 return false;
59         }
60     }
61     return true;
62 }

```

```

57         return false; // Matrices are not equal
58     }
59 }
60 }
61
62     return true; // Matrices are equal
63 }
64
65 /**
66  * Method to read a matrix from a Matrix Market file
67  * @param filename Name of the file
68  * @return SparseMatrix Sparse matrix
69  */
70 SparseMatrix readMatrixMarketFile(const std::string &filename)
71 {
72     std::ifstream file(filename); // Input file stream
73
74     // Check if the file was opened successfully
75     if (!file.is_open())
76     {
77         throw std::runtime_error("Unable to open file: " + filename);
78     }
79
80     std::string line; // String to hold the current line
81     bool isSymmetric = false, isPattern = false; // Flags to indicate if the matrix
82     // is symmetric or pattern only
83
84     // Skip the comments
85     while (std::getline(file, line))
86     {
87         // Check if the line is a comment
88         if (line[0] == '%')
89         {
90             // Check if the line contains the word "symmetric"
91             if (line.find("symmetric") != std::string::npos)
92             {
93                 isSymmetric = true; // Set the symmetric flag
94             }
95
96             // Check if the line contains the word "pattern"
97             if (line.find("pattern") != std::string::npos)
98             {
99                 isPattern = true; // Set the pattern flag
100             }
101         }
102         else
103         {
104             break; // First non-comment line reached, break out of the loop
105         }
106     }
107
108     // Read the matrix dimensions
109     int numRows, numCols, nonZeros; // Number of rows,
110     // columns and non-zero elements in the matrix
111     std::stringstream(line) >> numRows >> numCols >> nonZeros; // Read the
112     // dimensions from the line
113
114     // Check if the file was read successfully
115     if (!file)
116     {
117         throw std::runtime_error("Failed to read matrix dimensions from file: " +
118             filename);
119     }
120
121     SparseMatrix matrix; // Sparse
122     // matrix to hold the data
123     matrix.rowPtr.resize(numRows + 1, 0); // Resize
124     // the row pointer vector
125     std::vector<std::vector<std::pair<int, double>>> tempRows(numRows); //
126     // Temporary vector to hold the data

```

```

120     int rowIndex, colIndex;                                // Row and
        column indices
121     double value;                                          // Value of
        the non-zero element
122
123     // Read the non-zero elements
124     for (int i = 0; i < nonZeros; ++i)
125     {
126         // If the matrix is pattern only, the value of the non-zero element is 1.0
127         if (isPattern)
128         {
129             file >> rowIndex >> colIndex; // Read the row and column indices
130             value = 1.0;                  // Default value for pattern entries
131         }
132         else
133         {
134             file >> rowIndex >> colIndex >> value; // Read the row and column
                indices and the value
135         }
136
137         // Check if the file was read successfully
138         if (!file)
139         {
140             throw std::runtime_error("Failed to read data from file: " + filename);
141         }
142
143         rowIndex--; // Adjusting from 1-based to 0-based indexing
144         colIndex--; // Adjusting from 1-based to 0-based indexing
145
146         tempRows[rowIndex].emplace_back(colIndex, value); // Store the data in the
                temporary vector
147
148         // If the matrix is symmetric, store the data in the transpose as well
149         if (isSymmetric && rowIndex != colIndex)
150         {
151             tempRows[colIndex].emplace_back(rowIndex, value); // Store the data in
                the temporary vector
152         }
153     }
154
155     // Sort each row by column index
156     for (auto &row : tempRows)
157     {
158         std::sort(row.begin(), row.end());
159     }
160
161     // Reconstruct SparseMatrix structure
162     int cumSum = 0; // Cumulative sum of the number of non-zero elements
163
164     // Iterate over the rows of the matrix
165     for (int i = 0; i < numRows; ++i)
166     {
167         matrix.rowPtr[i] = cumSum; // Store the cumulative sum in the row pointer
                vector
168
169         // Iterate over the non-zero elements in the current row
170         for (const auto &elem : tempRows[i])
171         {
172             matrix.values.push_back(elem.second); // Store the value of the non-
                zero element
173             matrix.colIndices.push_back(elem.first); // Store the column index of
                the non-zero element
174         }
175
176         cumSum += tempRows[i].size(); // Update the cumulative sum
177     }
178
179     matrix.rowPtr[numRows] = cumSum; // Store the cumulative sum in the row pointer
        vector
180     matrix.numRows = numRows;        // Store the number of rows

```



```

181     matrix.numCols = numCols;           // Store the number of columns
182
183     // Return the sparse matrix
184     return matrix;
185 }
186
187 /**
188  * Method to generate a random fat vector
189  * @param n Number of rows
190  * @param m Number of columns
191  * @return FatVector fat vector
192  */
193 FatVector generateLargeFatVector(int n, int k)
194 {
195     FatVector denseVector(n, std::vector<double>(k)); // fat vector to hold the
196     random values
197
198     // Iterate over the rows of the fat vector
199     for (int i = 0; i < n; ++i)
200     {
201         // Iterate over the columns of the fat vector
202         for (int j = 0; j < k; ++j)
203         {
204             denseVector[i][j] = rand() % 100 + 1; // Generate a random value
205             between 1 and 100
206         }
207     }
208
209     // Return the fat vector
210     return denseVector;
211 }
212
213 /**
214  * @brief Method to serialize a FatVector to a flat array
215  * @param denseVec fat vector to serialize
216  * @return std::vector<double> Flat array containing the serialized data
217  */
218 std::vector<double> serialize(const FatVector &denseVec)
219 {
220     std::vector<double> flat; // Flat array to hold the serialized data
221
222     // Iterate over the rows of the fat vector
223     for (const auto &vec : denseVec)
224     {
225         flat.insert(flat.end(), vec.begin(), vec.end()); // Copy the elements
226     }
227
228     // Return the flat array
229     return flat;
230 }
231
232 /**
233  * @brief Method to deserialize a flat array to a FatVector
234  * @param flat Flat array to deserialize
235  * @param rows Number of rows in the fat vector
236  * @param cols Number of columns in the fat vector
237  * @return FatVector fat vector
238  */
239 FatVector deserialize(const std::vector<double> &flat, int rows, int cols)
240 {
241     FatVector denseVec(rows, std::vector<double>(cols)); // fat vector to hold the
242     deserialized data
243
244     // Iterate over the rows of the fat vector
245     for (int i = 0; i < rows; ++i)
246     {
247         // Iterate over the columns of the fat vector
248         for (int j = 0; j < cols; ++j)
249         {
250             denseVec[i][j] = flat[i * cols + j]; // Copy the element

```

```

248     }
249 }
250
251 // Return the fat vector
252 return denseVec;
253 }

```

## Appendix B.G Main File

Main file for running the different algorithms and comparing their performance.

```

1  #include "utils.h" // Utility functions
2  #include "SparseMatrixFatVectorMultiply.h" // Sequential algorithm
3  #include "SparseMatrixFatVectorMultiplyRowWise.h" // Parallel algorithm (row
   -wise)
4  #include "SparseMatrixFatVectorMultiplyColumnWise.h" // Parallel algorithm (
   column-wise)
5  #include "SparseMatrixFatVectorMultiplyNonZeroElement.h" // Parallel algorithm (non
   -zero element)
6
7  int main(int argc, char *argv[])
8  {
9      //
10
11     // ===== INITIALISATION
12     // =====
13
14     // Initialise MPI and PETSc
15     MPI_Init(&argc, &argv);
16     PetscInitialize(&argc, &argv, NULL, NULL);
17
18     // Retrieve the rank and size of the world communicator
19     int worldRank, worldSize;
20     MPI_Comm_rank(PETSC_COMM_WORLD, &worldRank);
21     MPI_Comm_size(PETSC_COMM_WORLD, &worldSize);
22
23     // Check if the correct number of arguments is provided
24     if (argc != 3)
25     {
26         if (worldRank == 0)
27         {
28             std::cerr << "Usage: " << argv[0] << " <number of columns> <matrix file
               path>" << std::endl;
29         }
30         MPI_Abort(PETSC_COMM_WORLD, 1);
31     }
32
33     // Parse the command-line arguments
34     int k = std::atoi(argv[1]); // Convert the first argument to an integer
35     std::string filename = argv[2]; // The second argument is the filename
36
37     // Declare the sparse matrix and Fat Vector
38     SparseMatrix M;
39     FatVector v;
40
41     // Declare the result of the serial multiplication
42     FatVector resultSerial;
43
44     // Declare the data for broadcasting the sparse matrix and Fat Vector
45     std::vector<double> flatData;
46     int dataSize = 0;
47
48     // Declare the variables for timing the execution of the algorithms

```

```

48     double startTime, endTime;
49
50     //
51     // ===== READ THE SPARSE MATRIX AND GENERATE THE Fat
52     // Vector =====
53
54     if (worldRank == 0)
55     {
56         std::cout << "World size: " << worldSize << std::endl;    // Print the
57         // number of processes
58         std::cout << "Sparse matrix: " << filename << std::endl; // Print the name
59         // of the Matrix Market file
60
61         // Read the sparse matrix from the Matrix Market file
62         M = readMatrixMarketFile(filename);
63         std::cout << "Matrix size: " << M.numRows << "x" << M.numCols << std::endl;
64
65         // Generate a random Fat Vector
66         v = generateLargeFatVector(M.numCols, k);
67         std::cout << "Vector size: " << M.numCols << "x" << k << std::endl;
68
69         // Prepare the data for broadcasting
70         flatData = serialize(v);    // Serialize the Fat Vector
71         dataSize = flatData.size(); // Size of the serialized data
72     }
73
74     // ===== EXECUTE THE SERIAL MULTIPLICATION =====
75
76     if (worldRank == 0)
77     {
78         // Execute the serial multiplication
79         startTime = MPI_Wtime();
80         resultSerial = sparseMatrixFatVectorMultiply(M, v, k);
81         endTime = MPI_Wtime();
82         std::cout << "Serial Algo Execution time: " << (endTime - startTime)
83         << std::endl;
84
85         // FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS OF THE RESULT
86         // std::cout << "Result: " << std::endl;
87         // for (int i = 0; i < 10; ++i)
88         // {
89         //     for (int j = 0; j < k; ++j)
90         //     {
91         //         std::cout << resultSerial[i][j] << " ";
92         //     }
93         //     std::cout << std::endl;
94         // }
95     }
96
97     // ===== BROADCAST THE SPARSE MATRIX AND Fat Vector =====
98
99     // Wait for the main process to finish the serial multiplication
100    MPI_Barrier(MPI_COMM_WORLD);

```

```

101 // ===== FOR DEBUGGING ONLY - START BROADCAST TIMER
102 // =====
103 // startTime = MPI_Wtime();
104 // ===== FOR DEBUGGING ONLY - START BROADCAST TIMER
105 // =====
106 // Broadcast the Sparse Matrix to all processes
107 // Prepare the data for broadcasting
108 int valuesSize = M.values.size(); // Number of non-
109     zero elements
110 int colIndicesSize = M.colIndices.size(); // Number of column
111     indices
112 int rowPtrSize = M.rowPtr.size(); // Number of row
113     pointers
114 MPI_Bcast(&M.numRows, 1, MPI_INT, 0, MPI_COMM_WORLD); // Broadcast the
115     number of rows
116 MPI_Bcast(&M.numCols, 1, MPI_INT, 0, MPI_COMM_WORLD); // Broadcast the
117     number of columns
118 MPI_Bcast(&valuesSize, 1, MPI_INT, 0, MPI_COMM_WORLD); // Broadcast the
119     number of non-zero elements
120 MPI_Bcast(&colIndicesSize, 1, MPI_INT, 0, MPI_COMM_WORLD); // Broadcast the
121     number of column indices
122 MPI_Bcast(&rowPtrSize, 1, MPI_INT, 0, MPI_COMM_WORLD); // Broadcast the
123     number of row pointers
124 // Resize the vectors for all processes
125 if (worldRank != 0)
126 {
127     M.values.resize(valuesSize);
128     M.colIndices.resize(colIndicesSize);
129     M.rowPtr.resize(rowPtrSize);
130 }
131 // Broadcast the data
132 MPI_Bcast(M.values.data(), valuesSize, MPI_DOUBLE, 0, MPI_COMM_WORLD);
133 MPI_Bcast(M.colIndices.data(), colIndicesSize, MPI_INT, 0, MPI_COMM_WORLD);
134 MPI_Bcast(M.rowPtr.data(), rowPtrSize, MPI_INT, 0, MPI_COMM_WORLD);
135
136 // Broadcast the Fat Vector to all processes
137 // Broadcast the size of the serialized data
138 MPI_Bcast(&dataSize, 1, MPI_INT, 0, MPI_COMM_WORLD);
139 // Resize flatData for all processes
140 if (worldRank != 0)
141 {
142     flatData.resize(dataSize);
143 }
144 // Broadcast the data
145 MPI_Bcast(flatData.data(), dataSize, MPI_DOUBLE, 0, MPI_COMM_WORLD);
146 // Deserialize the data
147 if (worldRank != 0)
148 {
149     v.resize(M.numCols, std::vector<double>(k));
150     v = deserialize(flatData, M.numCols, k);
151 }
152
153 // Wait for all processes to finish the broadcast
154 MPI_Barrier(MPI_COMM_WORLD);
155
156 // ===== FOR DEBUGGING ONLY - STOP BROADCAST TIMER
157 // =====
158 // endTime = MPI_Wtime();
159 // if (worldRank == 0)
160 // {
161 //     std::cout << "Broadcast time: " << (endTime - startTime) << std::endl;
162 // }
163 // ===== FOR DEBUGGING ONLY - STOP BROADCAST TIMER
164 // =====
165 //
166 // -----

```

```

157 // ===== EXECUTE THE PARALLEL MULTIPLICATION (ROW-WISE)
158 // -----
159
160 // Execute the parallel multiplication (row-wise)
161 startTime = MPI_Wtime();
162 FatVector resultRowWise = sparseMatrixFatVectorMultiplyRowWise(M, v, k);
163 endTime = MPI_Wtime();
164
165 // Only the main process prints the parallel execution time
166 if (worldRank == 0)
167 {
168     std::cout << "Row-wise Execution time: " << (endTime - startTime)
169     << std::endl;
170
171     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
172     // =====
173     // std::cout << "Result: " << std::endl;
174     // for (int i = 0; i < 10; ++i)
175     // {
176     //     for (int j = 0; j < k; ++j)
177     //     {
178     //         std::cout << resultRowWise[i][j] << " ";
179     //     }
180     //     std::cout << std::endl;
181     // }
182     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
183     // =====
184
185     // Compare the results of the serial and parallel multiplications
186     if (areMatricesEqual(resultSerial, resultRowWise, 1e-6)) // Tolerance = 1e
187     -6
188     {
189         std::cout << "Row-wise: Results are the same!"
190         << std::endl;
191     }
192     else
193     {
194         std::cout << "Row-wise: Results are different!"
195         << std::endl;
196     }
197 }
198
199 // -----
200
201 // ===== EXECUTE THE PARALLEL MULTIPLICATION (COLUMN-WISE)
202 // =====
203
204 // Wait for all processes to finish the parallel multiplication (row-wise)
205 MPI_Barrier(MPI_COMM_WORLD);
206
207 // Execute the parallel multiplication (column-wise)
208 startTime = MPI_Wtime();
209 FatVector resultColumnWise = sparseMatrixFatVectorMultiplyColumnWise(M, v, k);
210 endTime = MPI_Wtime();
211
212 // Only the main process prints the parallel execution time
213 if (worldRank == 0)
214 {
215     std::cout << "Column-wise Execution time: " << (endTime - startTime)
216     << std::endl;
217
218     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
219     // =====

```

```

215 // std::cout << "Result: " << std::endl;
216 // for (int i = 0; i < 10; ++i)
217 // {
218 //     for (int j = 0; j < k; ++j)
219 //     {
220 //         std::cout << resultColumnWise[i][j] << " ";
221 //     }
222 //     std::cout << std::endl;
223 // }
224 // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
225 // =====
226
227 // Compare the results of the serial and parallel multiplications
228 if (areMatricesEqual(resultSerial, resultColumnWise, 1e-6)) // Tolerance =
229     1e-6
230 {
231     std::cout << "Column-wise: Results are the same!"
232     << std::endl;
233 }
234 else
235 {
236     std::cout << "Column-wise: Results are different!"
237     << std::endl;
238 }
239 //
240 // -----
241
242 // ===== EXECUTE THE PARALLEL MULTIPLICATION (NON-ZERO
243 // ELEMENT) =====
244 // -----
245
246 // Wait for all processes to finish the parallel multiplication (column-wise)
247 MPI_Barrier(MPI_COMM_WORLD);
248
249 // Execute the parallel multiplication (non-zero element)
250 startTime = MPI_Wtime();
251 FatVector resultNonZeroElement = sparseMatrixFatVectorMultiplyNonZeroElement(M,
252     v, k);
253 endTime = MPI_Wtime();
254
255 // Only the main process prints the parallel execution time
256 if (worldRank == 0)
257 {
258     std::cout << "Non-zero Elements Execution time: " << (endTime - startTime)
259     << std::endl;
260
261     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
262     // =====
263     // std::cout << "Result: " << std::endl;
264     // for (int i = 0; i < 10; ++i)
265     // {
266     //     for (int j = 0; j < k; ++j)
267     //     {
268     //         std::cout << resultNonZeroElement[i][j] << " ";
269     //     }
270     //     std::cout << std::endl;
271     // }
272     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
273     // =====
274
275     // Compare the results of the serial and parallel multiplications
276 if (areMatricesEqual(resultSerial, resultNonZeroElement, 1e-6)) //
277     Tolerance = 1e-6
278 {
279     std::cout << "Non-zero Elements: Results are the same!"
280     << std::endl;

```

```

274     }
275     else
276     {
277         std::cout << "Non-zero Elements: Results are different!"
278                 << std::endl;
279     }
280 }
281
282 //
283 // ===== EXECUTE THE PARALLEL MULTIPLICATION (
284 // PETSc) =====
285
286 // Wait for all processes to finish the parallel multiplication (non-zero
287 // element)
288 MPI_Barrier(MPI_COMM_WORLD);
289
290 // Declare the PETSc matrix
291 Mat A, B, C;
292
293 // ===== FOR DEBUGGING ONLY - START PETSCS SETUP TIMER
294 // =====
295 // startTime = MPI_Wtime();
296 // ===== FOR DEBUGGING ONLY - START PETSCS SETUP TIMER
297 // =====
298
299 // Create a parallel matrix to store the sparse matrix
300 MatCreate(PETSC_COMM_WORLD, &A);
301 MatSetSizes(A, PETSC_DECIDE, PETSC_DECIDE, M.numRows, M.numCols);
302 MatSetType(A, MATMPIAIJ);
303 MatSetUp(A);
304 // Fill the PETSc matrix with the values from the sparse matrix
305 if (worldRank == 0)
306 {
307     for (int i = 0; i < M.numRows; ++i)
308     {
309         for (int j = M.rowPtr[i]; j < M.rowPtr[i + 1]; ++j)
310         {
311             MatSetValue(A, i, M.colIndices[j], M.values[j], INSERT_VALUES);
312         }
313     }
314 }
315 // Assemble the PETSc matrix
316 MatAssemblyBegin(A, MAT_FINAL_ASSEMBLY);
317 MatAssemblyEnd(A, MAT_FINAL_ASSEMBLY);
318
319 // Create a parallel matrix to store the Fat Vector
320 MatCreate(PETSC_COMM_WORLD, &B);
321 MatSetSizes(B, PETSC_DECIDE, PETSC_DECIDE, M.numCols, k);
322 MatSetType(B, MATDENSE);
323 MatSetUp(B);
324 // Fill the PETSc matrix B with values from the dense matrix v
325 if (worldRank == 0)
326 {
327     for (int i = 0; i < M.numCols; ++i)
328     {
329         for (int j = 0; j < k; ++j)
330         {
331             MatSetValue(B, i, j, v[i][j], INSERT_VALUES);
332         }
333     }
334 }
335 // Assemble the PETSc matrix
336 MatAssemblyBegin(B, MAT_FINAL_ASSEMBLY);
337 MatAssemblyEnd(B, MAT_FINAL_ASSEMBLY);

```

```

336 // ===== FOR DEBUGGING ONLY - STOP PETSCS SETUP TIMER
337 // =====
338 // endTime = MPI_Wtime();
339 // if (worldRank == 0)
340 // {
341 //     std::cout << "PETSc Setup time: " << (endTime - startTime) << std::endl;
342 // }
343 // ===== FOR DEBUGGING ONLY - STOP PETSCS SETUP TIMER
344 // =====
345 // Create a parallel matrix to store the result of the multiplication
346 startTime = MPI_Wtime();
347 MatProductCreate(A, B, NULL, &C);
348 MatMatMult(A, B, MAT_INITIAL_MATRIX, PETSC_DEFAULT, &C);
349 endTime = MPI_Wtime();
350 if (worldRank == 0)
351 {
352     // Print the execution time
353     std::cout << "PETSc Execution time: " << (endTime - startTime) << std::endl;
354 }
355 // ===== FOR DEBUGGING ONLY - START PETSCS CONVERSION
356 // TIMER =====
357 // startTime = MPI_Wtime();
358 // ===== FOR DEBUGGING ONLY - START PETSCS CONVERSION
359 // TIMER =====
360 // Create a sequential matrix to retrieve the result
361 Mat CSeq;
362 MatCreateRedundantMatrix(C, worldSize, MPI_COMM_NULL, MAT_INITIAL_MATRIX, &CSeq);
363
364 if (worldRank == 0)
365 {
366     // Convert the result matrix C to a FatVector
367     FatVector globalMatrix = ConvertPETScMatToFatVector(CSeq);
368
369     // ===== FOR DEBUGGING ONLY - STOP PETSCS CONVERSION
370     // TIMER =====
371     // endTime = MPI_Wtime();
372     // std::cout << "PETSc Conversion time: " << (endTime - startTime) << std::endl;
373     // ===== FOR DEBUGGING ONLY - STOP PETSCS CONVERSION
374     // TIMER =====
375
376     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
377     // =====
378     // std::cout << "Result: " << std::endl;
379     // for (int i = 0; i < 10; ++i)
380     // {
381     //     for (int j = 0; j < k; ++j)
382     //     {
383     //         std::cout << globalMatrix[i][j] << " ";
384     //     }
385     //     std::cout << std::endl;
386     // }
387     // ===== FOR DEBUGGING ONLY - PRINT 10 FIRST ELEMENTS
388     // =====
389
390     // Compare the results of the serial and PETSc multiplications
391     if (areMatricesEqual(resultSerial, globalMatrix, 1e-6)) // Tolerance = 1e-6
392     {
393         std::cout << "PETSc: Results are the same!"
394         << std::endl;
395     }
396     else
397     {
398         std::cout << "PETSc: Results are different!"
399         << std::endl;
400     }
401 }

```



```

395     }
396 }
397
398 // Free the memory
399 MatDestroy(&A);
400 MatDestroy(&B);
401 MatDestroy(&C);
402 MatDestroy(&CSeq);
403
404 // Finalise MPI and PETSc
405 PetscFinalize();
406 MPI_Finalize();
407
408 return 0;
409 }

```

## Appendix B.H Scripts

### B.H.1 MPI Submission Script

Bash script to submit an MPI job to the cluster.

```

#!/bin/bash
##
## MPI submission script for PBS on CR2
## -----
##
## "MPI-sub2022v1"
## Follow the 6 steps below to configure your job
##
## STEP 1:
##
## Enter a job name after the -N on the line below:
##
##PBS -N mpi_assessment_test_8_cores_16_425500
##
## STEP 2:
##
## Select the number of cpus/cores required by modifying the #PBS -l select line
## below
##
## Normally you select cpus in chunks of 16 cpus
## The Maximum value for ncpus is 16 and mpirprocs MUST be the same value as ncpus.
##
## If more than 16 cpus are required then select multiple chunks of 16
## e.g. 16 CPUs: select=1:ncpus=16:mpiprocs=16
##      32 CPUs: select=2:ncpus=16:mpiprocs=16
##      ..etc..
##
##PBS -l select=2:ncpus=16:mpiprocs=16
##
## STEP 3:
##
## Select the correct queue by modifying the #PBS -q line below
##
## half_hour      - 30 minutes
## one_hour       - 1 hour
## three_hour     - 3 hours
## six_hour       - 6 hours
## half_day       - 12 hours
## one_day        - 24 hours
## two_day        - 48 hours
## five_day       - 120 hours
## ten_day        - 240 hours (by special arrangement)
##
##PBS -q half_hour

```

```

##
## STEP 4:
##
## Replace the hpc@cranfield.ac.uk email address
## with your Cranfield email address on the #PBS -M line below:
## Your email address is NOT your username
##
#PBS -m abe
#PBS -M alexis.balayre.500@cranfield.ac.uk
##
## =====
## DO NOT CHANGE THE LINES BETWEEN HERE
## =====
#PBS -j oe
#PBS -W sandbox=PRIVATE
#PBS -k n
ln -s $PWD $PBS_O_WORKDIR/$PBS_JOBID
## Change to working directory
cd $PBS_O_WORKDIR
## Calculate number of CPUs
export cpus='cat $PBS_NODEFILE | wc -l'
sort -u $PBS_NODEFILE -o mpi_nodes.$$
export I_MPI_HYDRA_IFACE=ib0
export I_MPI_HYDRA_BOOTSTRAP=ssh
export I_MPI_HYDRA_RMK=pbs
export K_VALUE=1
export MATRIX_PATH=/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/cop20k_A.
    mtx
## Debug options - only enable when instructed by HPC support
##export I_MPI_HYDRA_DEBUG=1
##export I_MPI_DEBUG=6
##export I_MPI_DEBUG_OUTPUT=%h-%r-%p-debug.out
## set some MPI tuning parameters to use the correct transport
## =====
## AND HERE
## =====
##
## STEP 5:
##
## Load the default application environment
## For a specific version add the version number, e.g.
## module load intel/2016b
##
module use /apps/modules/all
module load intel/2021b
##
## STEP 6:
##
## Run MPI code
##
## The main parameter to modify is your mpi program name
## - change YOUR_EXECUTABLE to your own filename
##

mpirun -genval1 -hostfile mpi_nodes.$$ -np ${cpus} ../my_program_final_debug ${
    K_VALUE} ${MATRIX_PATH}

## Tidy up the log directory
## DO NOT CHANGE THE LINE BELOW
## =====
rm $PBS_O_WORKDIR/$PBS_JOBID
#

```

## B.H.2 Batch Test Script

Bash script to submit multiple MPI jobs to the cluster.

```
#!/bin/bash

# Script to submit a batch of jobs to the cluster

# Path to the original script
original_script="mpi.sub"

# Maximum number of cores used for the job
max_cores=96

# Define a set of k values to test (Number of columns in the Fat Vector)
k_values=(1 3 6 9 12)

# Define a set of paths to test
paths=(
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/cop20k_A.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/adder_dcop_32.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/bcsstk17.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/af23560.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/amazon0302.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/cavity10.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/cage4.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/dc1.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/FEM_3D_thermal1.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/mac_econ_fwd500.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/mcfe.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/mhd4800a.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/olafu.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/raefsky2.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/rdist2.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/thermal1.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/thermomech_TK.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/west2021.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/lung2.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/olm1000.mtx"
    "/mnt/beegfs/home/s425500/hpc/assignment/sparse-matrix/roadNet-PA.mtx"
)

# Loop over the k values
for k_value in "${k_values[@]"; do
    # Loop over the paths of MTX files
    for path in "${paths[@]"; do
        # Loop over the number of chunks
        for chunks in $(seq 1 $((max_cores / 16))); do
            # Loop over the number of cpus per chunk
            for cpus in $(seq 2 16); do
                # Calculate the total number of cores
                total_cores=$((chunks * cpus))
                # Check if the total number of cores is less than the maximum
                number of cores
                if [ $total_cores -le $max_cores ]; then
                    echo "Submitting job with $total_cores cores, $chunks chunks
and $cpus cpus per chunk"
                    echo "Path: $path"

                    # Create a unique job name
                    matrix_name=$(basename "$path") # Remove the path
                    sanitized_matrix_name=${matrix_name//[a-zA-Z0-9_]/_} # Replace
all non-alphanumeric characters with underscores
                    job_name="${sanitized_matrix_name}_k${k_value}_cores${
total_cores}_chunks${chunks}_cpus${cpus}" # Add the k value to the job name

                    # Create a temporary submission script
                    temp_script="temp_${job_name}.sub"
                    cp "$original_script" "$temp_script"
```

```

# Replace the variables in the temporary script
sed -i "s|export k_value=.*|export k_value=${k_value}|" "
$temp_script" # Export the k value
sed -i "s|export MATRIX_PATH=.*|export MATRIX_PATH=${path}|" "
$temp_script" # Export the path to the MTX file
sed -i "s|#PBS -N .*|#PBS -N $job_name|" "$temp_script" # Set
the job name
sed -i "s|#PBS -l select=.*|#PBS -l select=${chunks}:ncpus=
$cpus:mpiprocs=$cpus|" "$temp_script" # Set the number of chunks and cpus per
chunk

# Submit the job and get the job id
job_id=$(qsub "$temp_script")
echo "Job id: $job_id"

# Wait until the job is finished
while true; do
    # Get the job status and duration
    job_status=$(qstat -f "$job_id" | grep job_state | awk '{
print $3}') # Get the job status
    job_duration=$(qstat -f "$job_id" | grep resources_used.
walltime | awk '{print $3}') # Get the job duration
    job_duration_seconds=$(echo $job_duration | awk -F: '{
print ($1 * 3600) + ($2 * 60) + $3 }') # Convert the job duration to seconds
    echo "Job status: $job_status"
    echo "Job duration: $job_duration"

    # If the job is finished, break the loop
    if [ -z "$job_status" ]; then
        break
    fi

    # if the job is running for more than 4 minutes, cancel it
    if [ "$job_duration_seconds" -gt 240 ]; then
        echo "Job is running for more than 4 minutes.
Cancelling it."
        qdel "$job_id"
        break
    fi

    # Wait for 1 second
    sleep 1
done

# Remove the temporary script
rm "$temp_script"
fi
done
done
done
done

```

### B.H.3 Get CSV Script

Bash script to analyse all job results files and extract the relevant information to create a CSV file.

```
#!/bin/bash

# Name of the CSV file to write the data to
output_csv="results.csv"

# Headers for the CSV file
echo "file Name,Cores Number,Sparse Matrix,Fat Vector,Serial Algo Execution time,
Row-wise Average Communication Time,Row-wise Average Computation Time,Row-wise
Execution time,Row-wise Result,Column-wise Average Communication Time,Column-
wise Average Computation Time,Column-wise Execution time,Column-wise Result,Non
-zero elements Average Communication Time,Non-zero elements Average Computation
Time,Non-zero Elements Execution time,Non-zero Elements Result,PETSc Execution
time,PETSc Result" >$output_csv

# Loop over the output files
for file in *.o*; do
    # Check that the file is valid and that it is a result file
    if [[ -s $file && $file == *mtx* ]]; then
        # Extract the job name and the number of cores from the file name
        job_name=$(basename "$file" | sed -e 's/\.[^.]*/$//') # Remove file
        extension
        num_cores=$(echo $file | grep -oP '(?<=_cores)\d+') # Extract the number
        of cores from the file name

        # Extract the matrix size and the vector size from the file
        matrix_size=$(grep "Matrix size" $file | awk '{print $3}' | sed 's/size://')
        ) # Extract the matrix size from the file
        vector_size=$(grep "Vector size" $file | awk '{print $3}' | sed 's/size://')
        ) # Extract the vector size from the file

        # Extract the serial execution time from the file
        serial_time=$(grep "Serial Algo Execution time" $file | awk '{print $5}')

        # Row-wise Data
        row_wise_communication_time=$(grep "Row-wise Average Communication Time"
$file | awk '{print $5}') # Extract the row-wise average communication time
from the file
        row_wise_computation_time=$(grep "Row-wise Average Computation Time" $file
| awk '{print $5}') # Extract the row-wise average computation time from
the file
        row_wise_execution_time=$(grep "Row-wise Execution time" $file | awk '{
print $4}') # Extract the row-wise execution time from the file
        row_wise_result=$(grep "Row-wise: Results are" $file | awk '{print $5}')
        # Extract the row-wise result from the file
        row_wise_result=$(if [ $row_wise_result == "same!" ]; then echo "same";
else echo "different"; fi) # Convert the row-wise result to a boolean

        # Column-wise Data
        col_wise_communication_time=$(grep "Column-wise Average Communication Time"
$file | awk '{print $6}') # Extract the column-wise average communication time
from the file
        col_wise_computation_time=$(grep "Column-wise Average Computation Time"
$file | awk '{print $6}') # Extract the column-wise average computation
time from the file
        col_wise_execution_time=$(grep "Column-wise Execution time" $file | awk '{
print $4}') # Extract the column-wise execution time from the
file
        col_wise_result=$(grep "Column-wise: Results are" $file | awk '{print $5}')
        # Extract the column-wise result from the file
        col_wise_result=$(if [ $col_wise_result == "same!" ]; then echo "same";
else echo "different"; fi) # Convert the column-wise result to a boolean

        # Non-zero element Data
        nonzero_communication_time=$(grep "Non-zero elements Average Communication
Time" $file | awk '{print $6}') # Extract the non-zero elements average
```

```

communication time from the file
    nonzero_computation_time=$(grep "Non-zero elements Average Computation Time
" $file | awk '{print $6}') # Extract the non-zero elements average
computation time from the file
    nonzero_execution_time=$(grep "Non-zero Elements Execution time" $file |
awk '{print $5}') # Extract the non-zero elements execution
time from the file
    nonzero_result=$(grep "Non-zero Elements: Results are" $file | awk '{print
$6}') # Extract the non-zero elements result from the
file
    nonzero_result=$(if [ $nonzero_result == "same!" ]; then echo "same"; else
echo "different"; fi) # Convert the non-zero elements result to a
boolean

# PETSc Data
    petsc_execution_time=$(grep "PETSc Execution time" $file | awk '{print $4}'
) # Extract the PETSc execution time from the file
    petsc_result=$(grep "PETSc: Results are" $file | awk '{print $5}')
# Extract the PETSc result from the file
    petsc_result=$(if [ $petsc_result == "same!" ]; then echo "same"; else echo
"different"; fi) # Convert the PETSc result to a boolean

# Write the extracted data to the CSV file
    echo "$job_name,$num_cores,$matrix_size,$vector_size,$serial_time,
$row_wise_communication_time,$row_wise_computation_time,
$row_wise_execution_time,$row_wise_result,$col_wise_communication_time,
$col_wise_computation_time,$col_wise_execution_time,$col_wise_result,
$nonzero_communication_time,$nonzero_computation_time,$nonzero_execution_time,
$nonzero_result,$petsc_execution_time,$petsc_result" >>$output_csv
fi
done

echo "The data was successfully written in $output_csv"

```